

# Convolutional Neural Networks and Image Processing

Jiří Materna



# About me

- Ph.D. in Natural Language Processing and Artificial Intelligence at Masaryk University
- 10 years at Seznam.cz (last 8 years as Head Of Research)
- Founder and lecturer at ML College
- Founder and co-organizer of ML Prague
- ML Freelance and consultant

# Outline

## Day 1 and 2

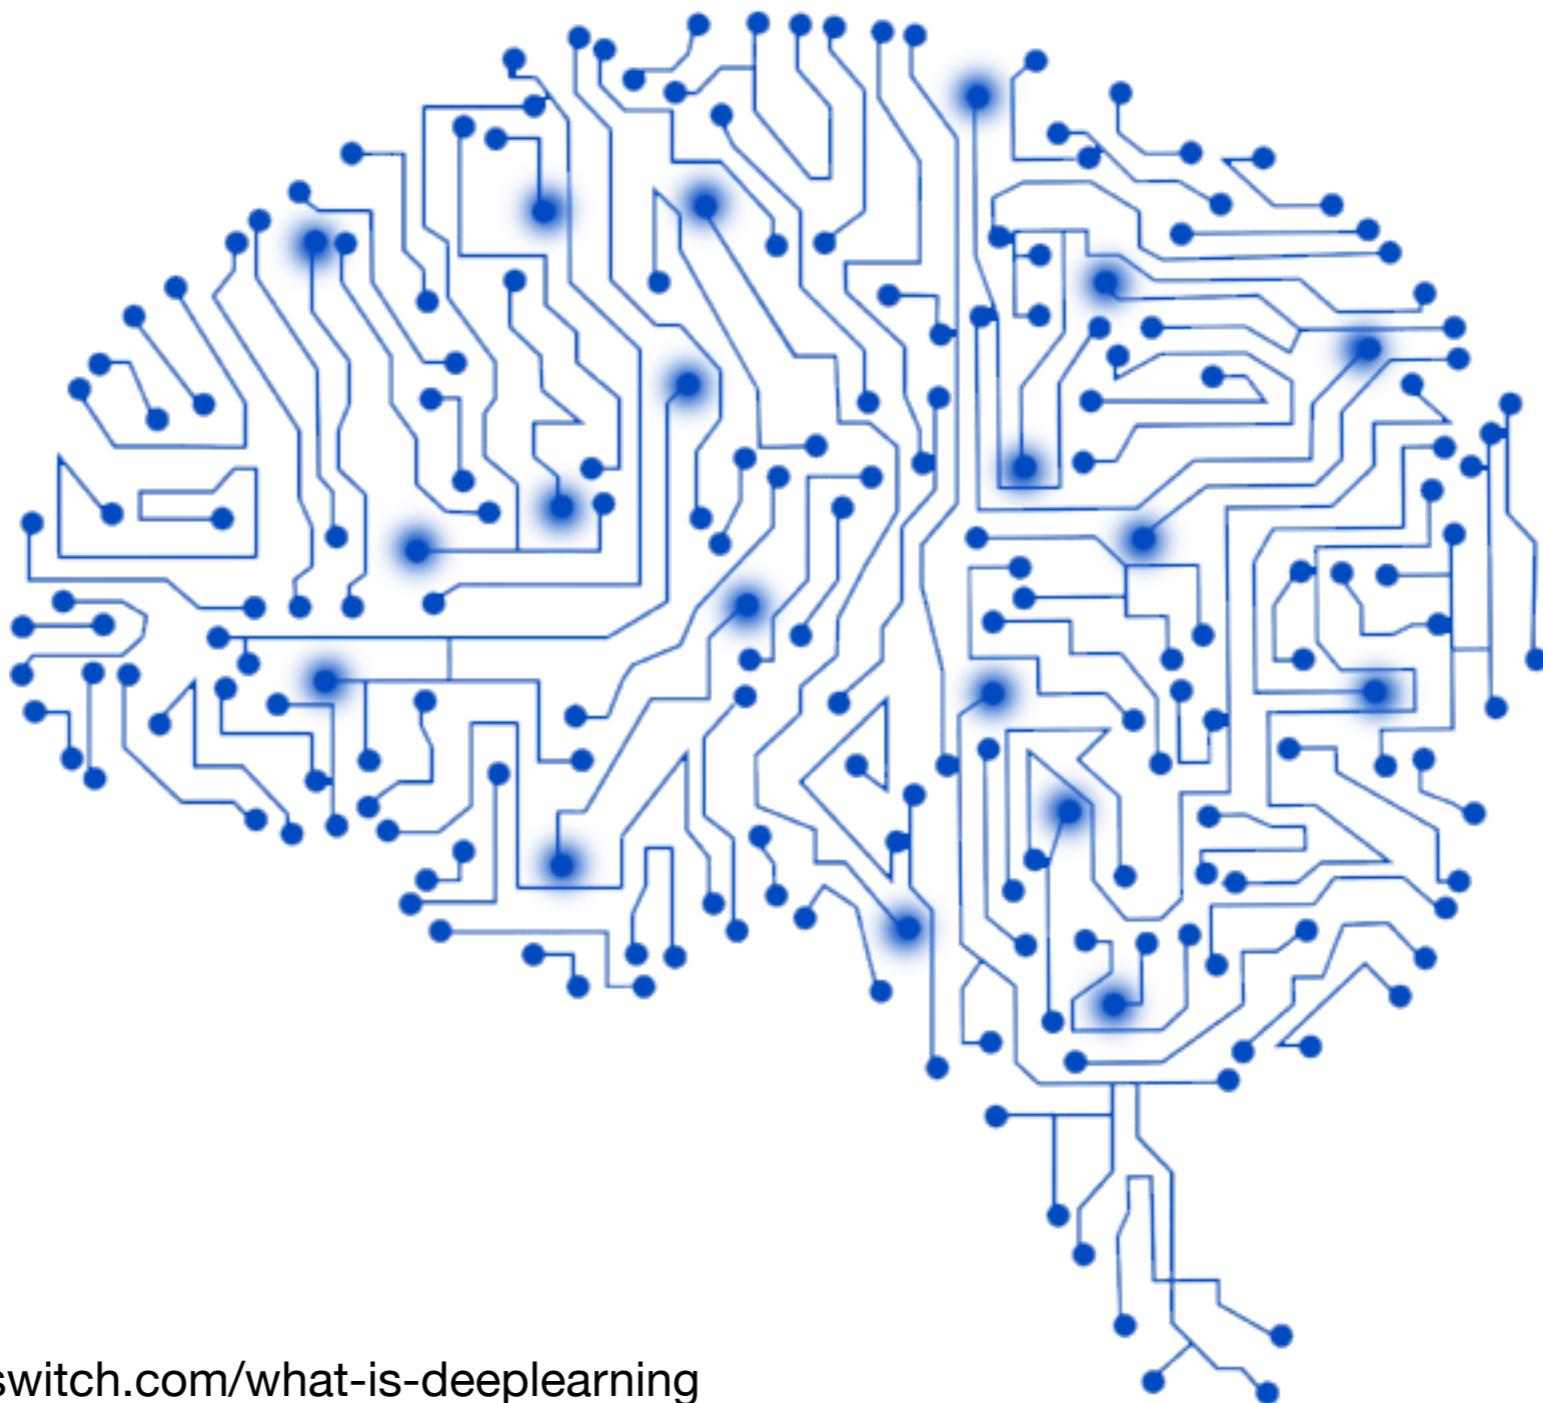
- Introduction to neural networks
- Activation functions for neural networks
- Multilayered neural networks
- Methods for training neural networks
- Convolutional neural networks
- Keras tutorial
- Practical classification and regression tasks solved using neural networks

# Outline

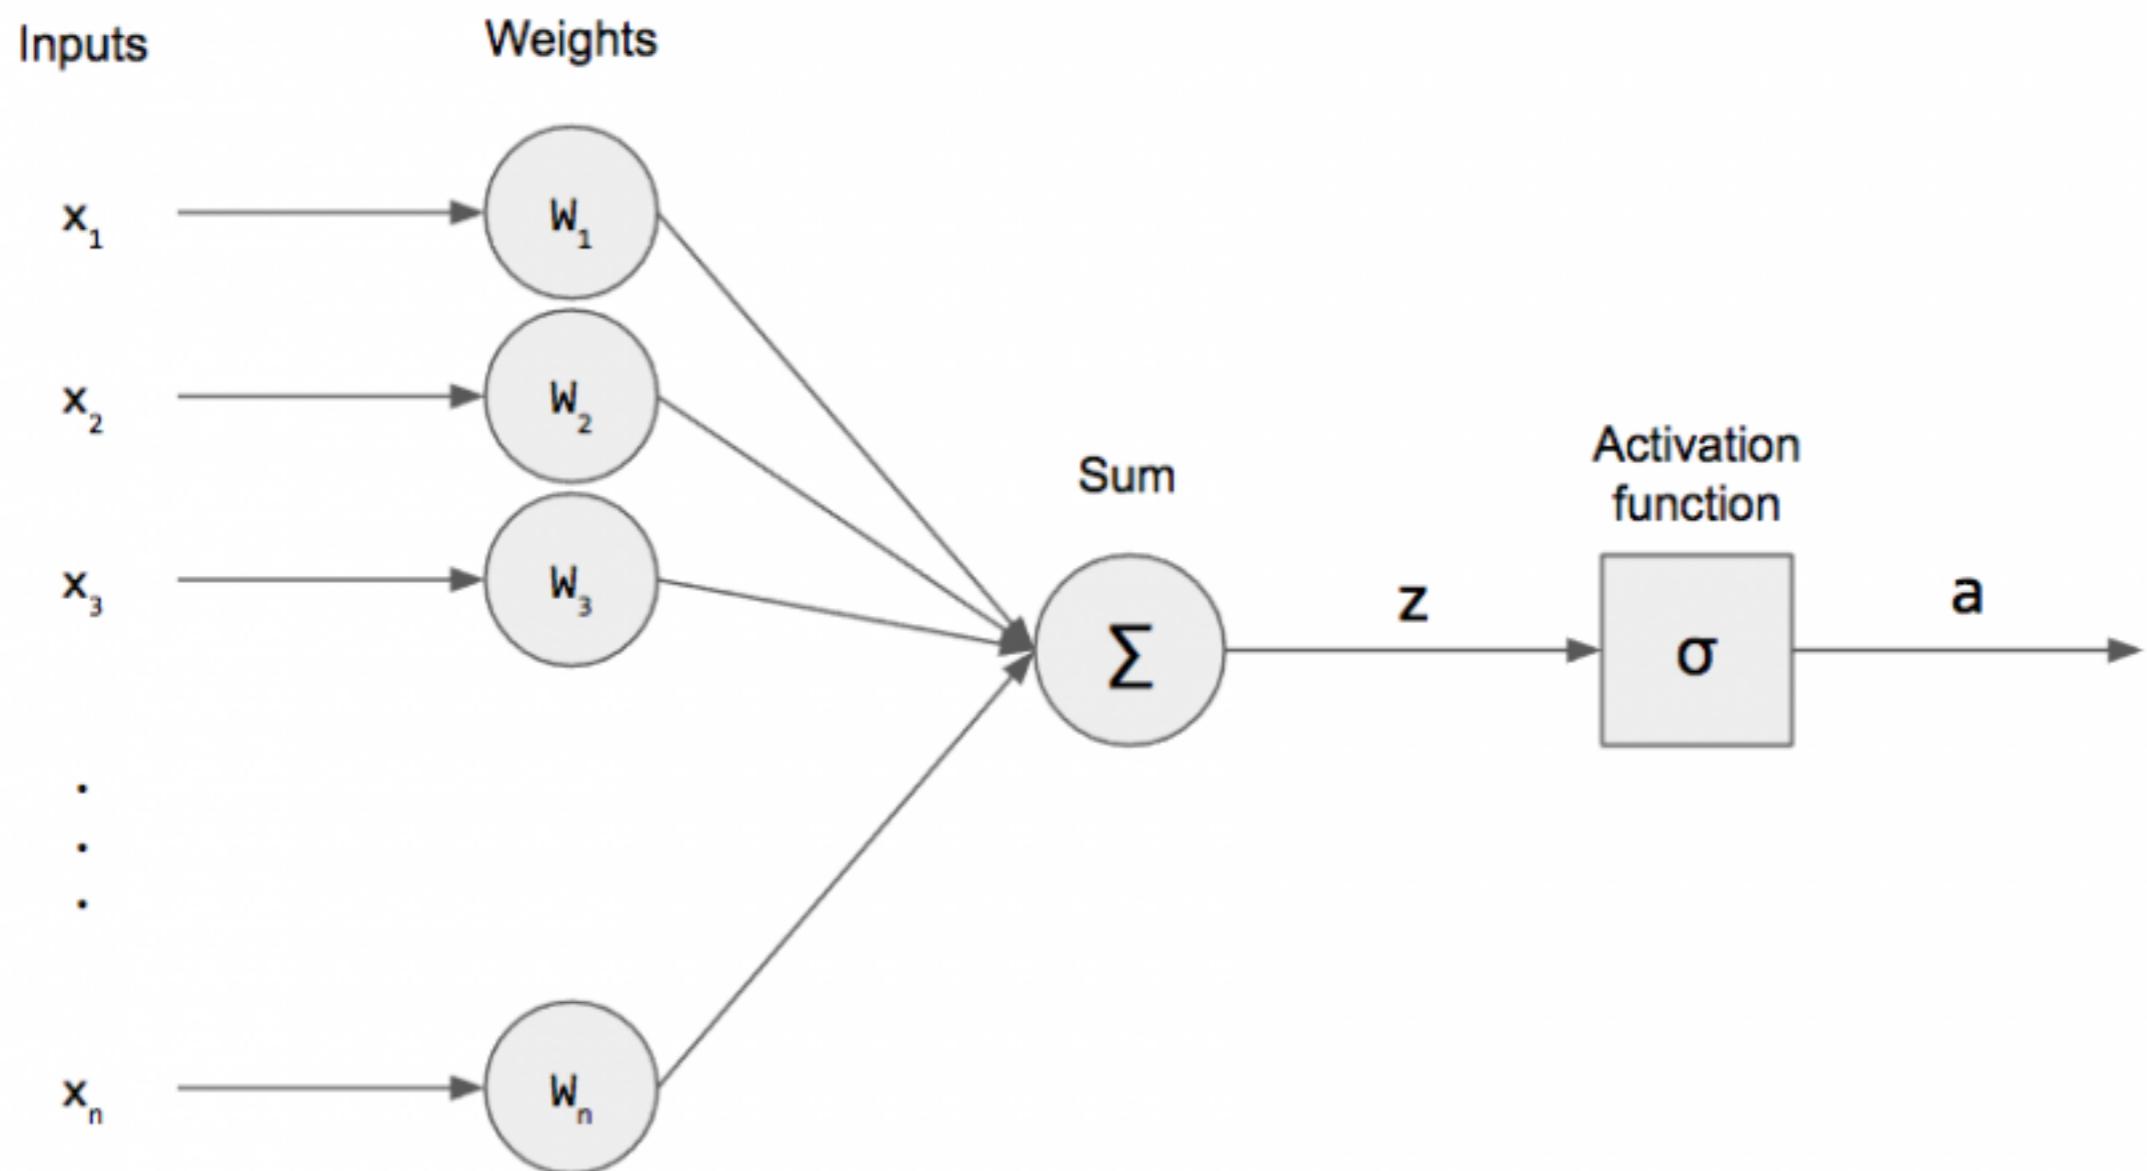
## Day 3 and 4

- VGG 16 and ResNet
- Transfer learning and fine-tuning
- Image classification
- Batch normalization and data augmentation
- U-net and Image segmentation
- GANs and superresolution
- Neural network explainability
- Adversarial patch

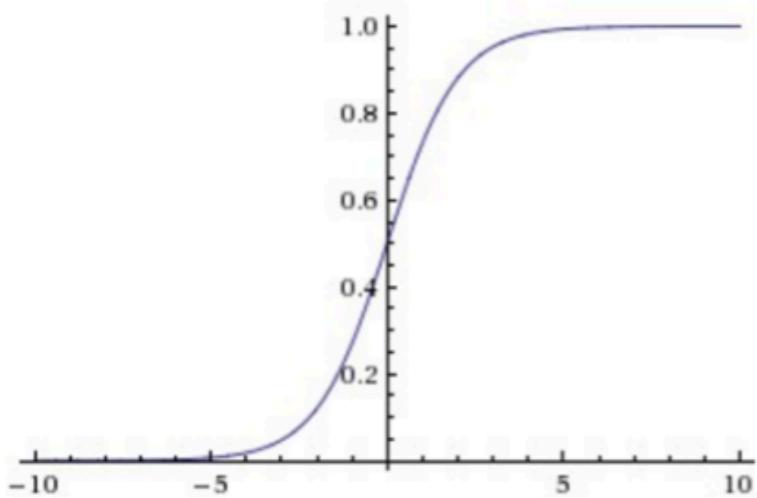
# Neural networks and deep learning



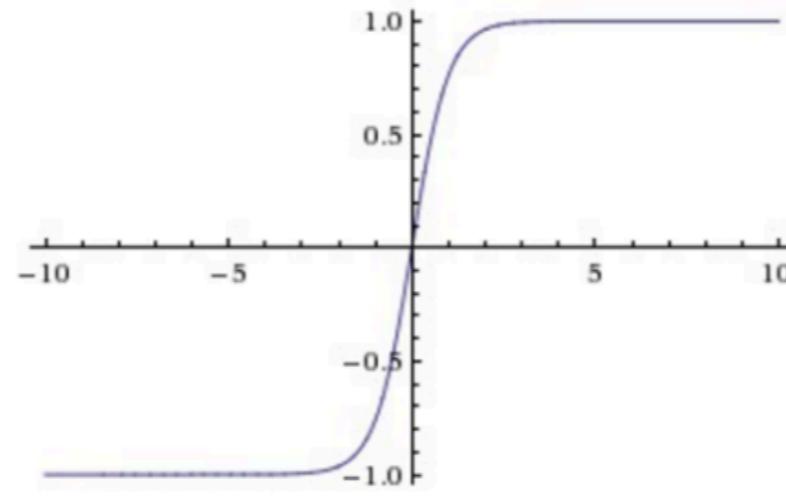
# Perceptron



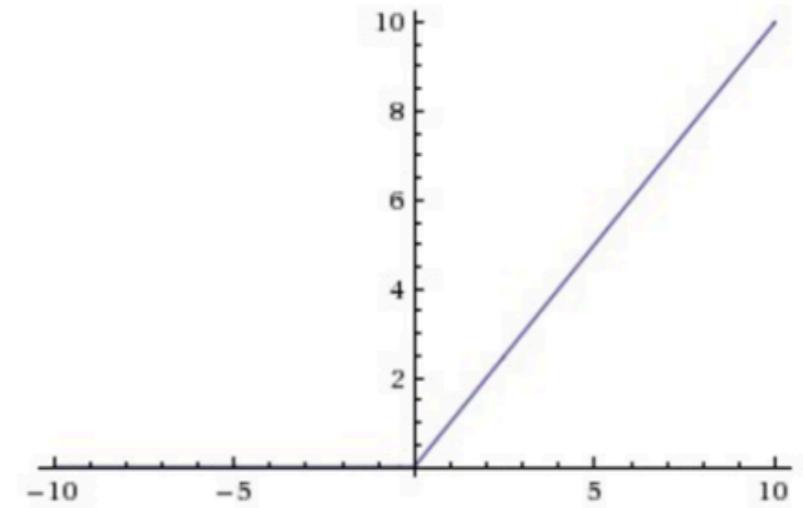
# Activation functions



Sigmoid



tanh

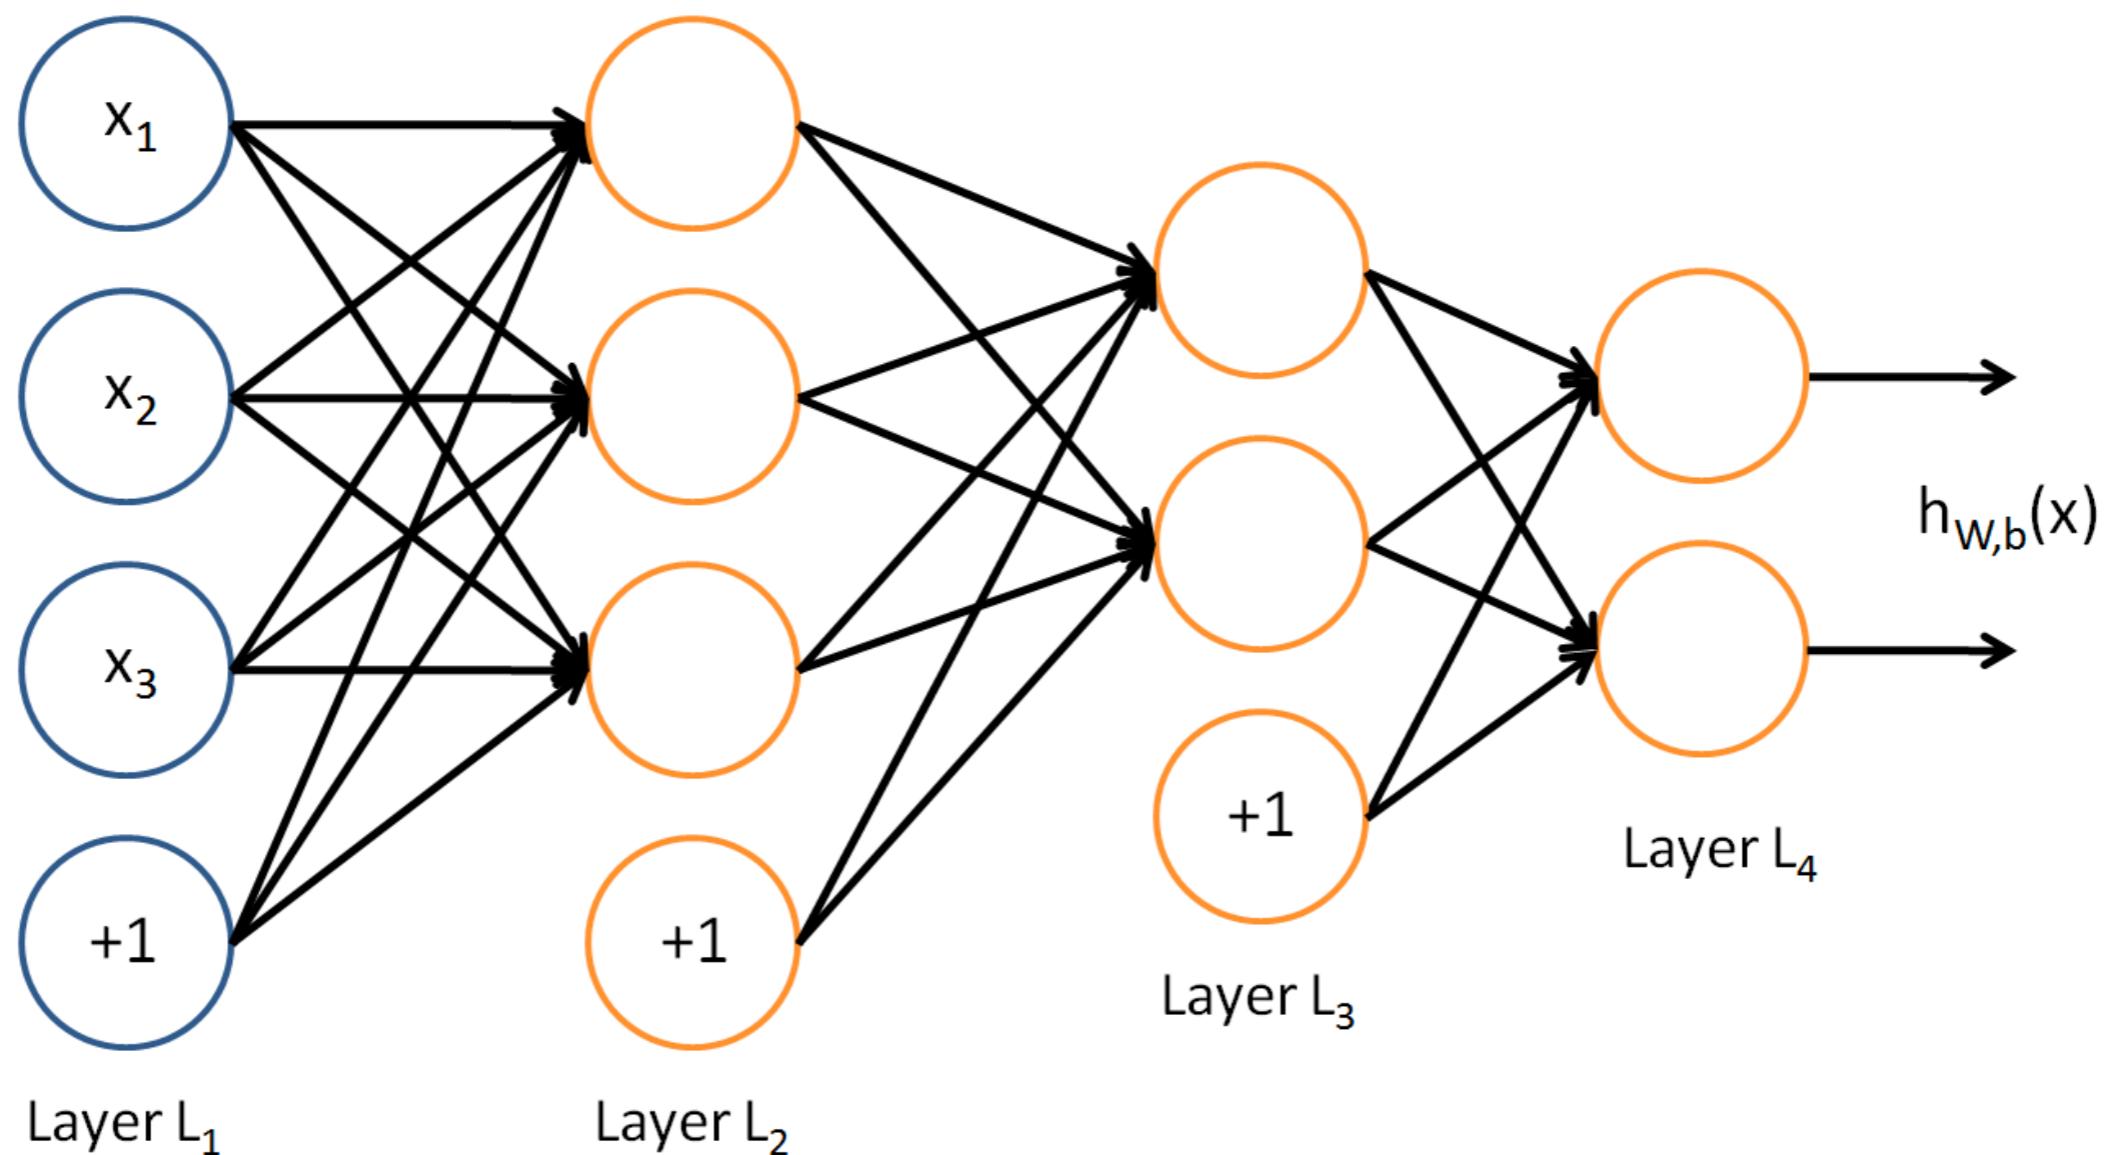


ReLU

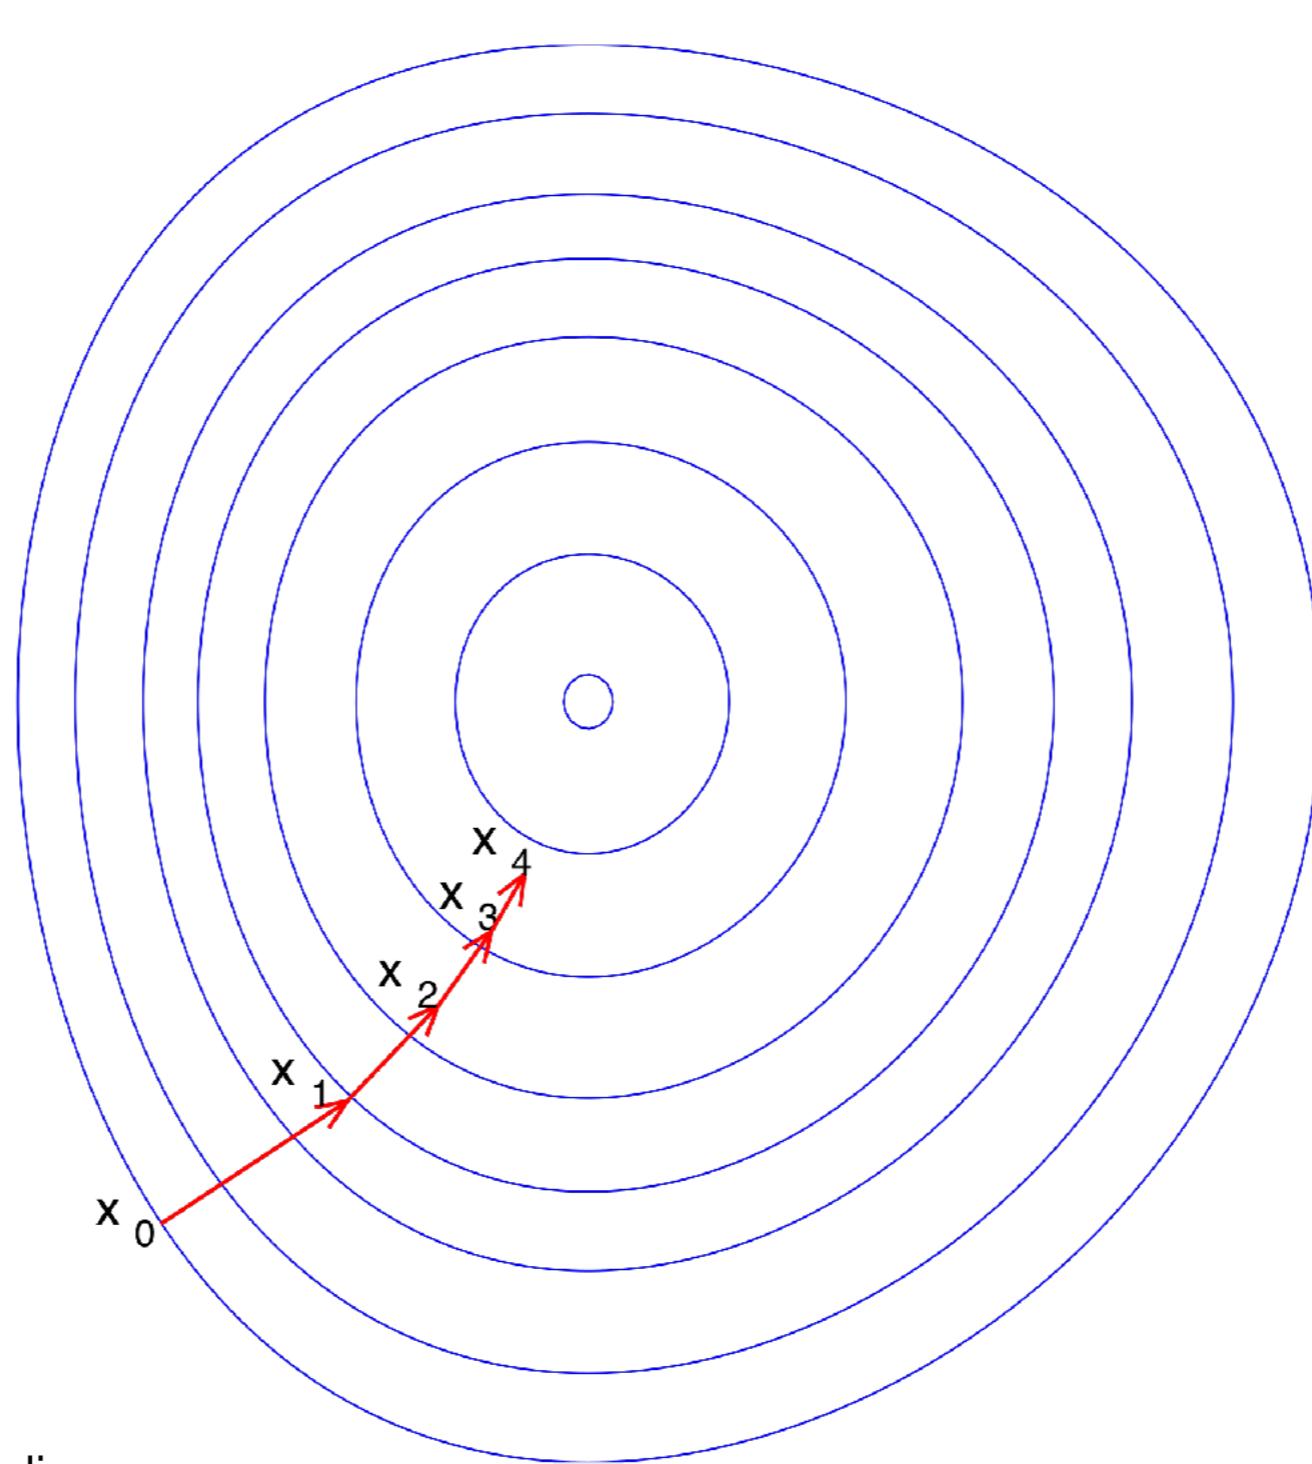
Softmax:

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

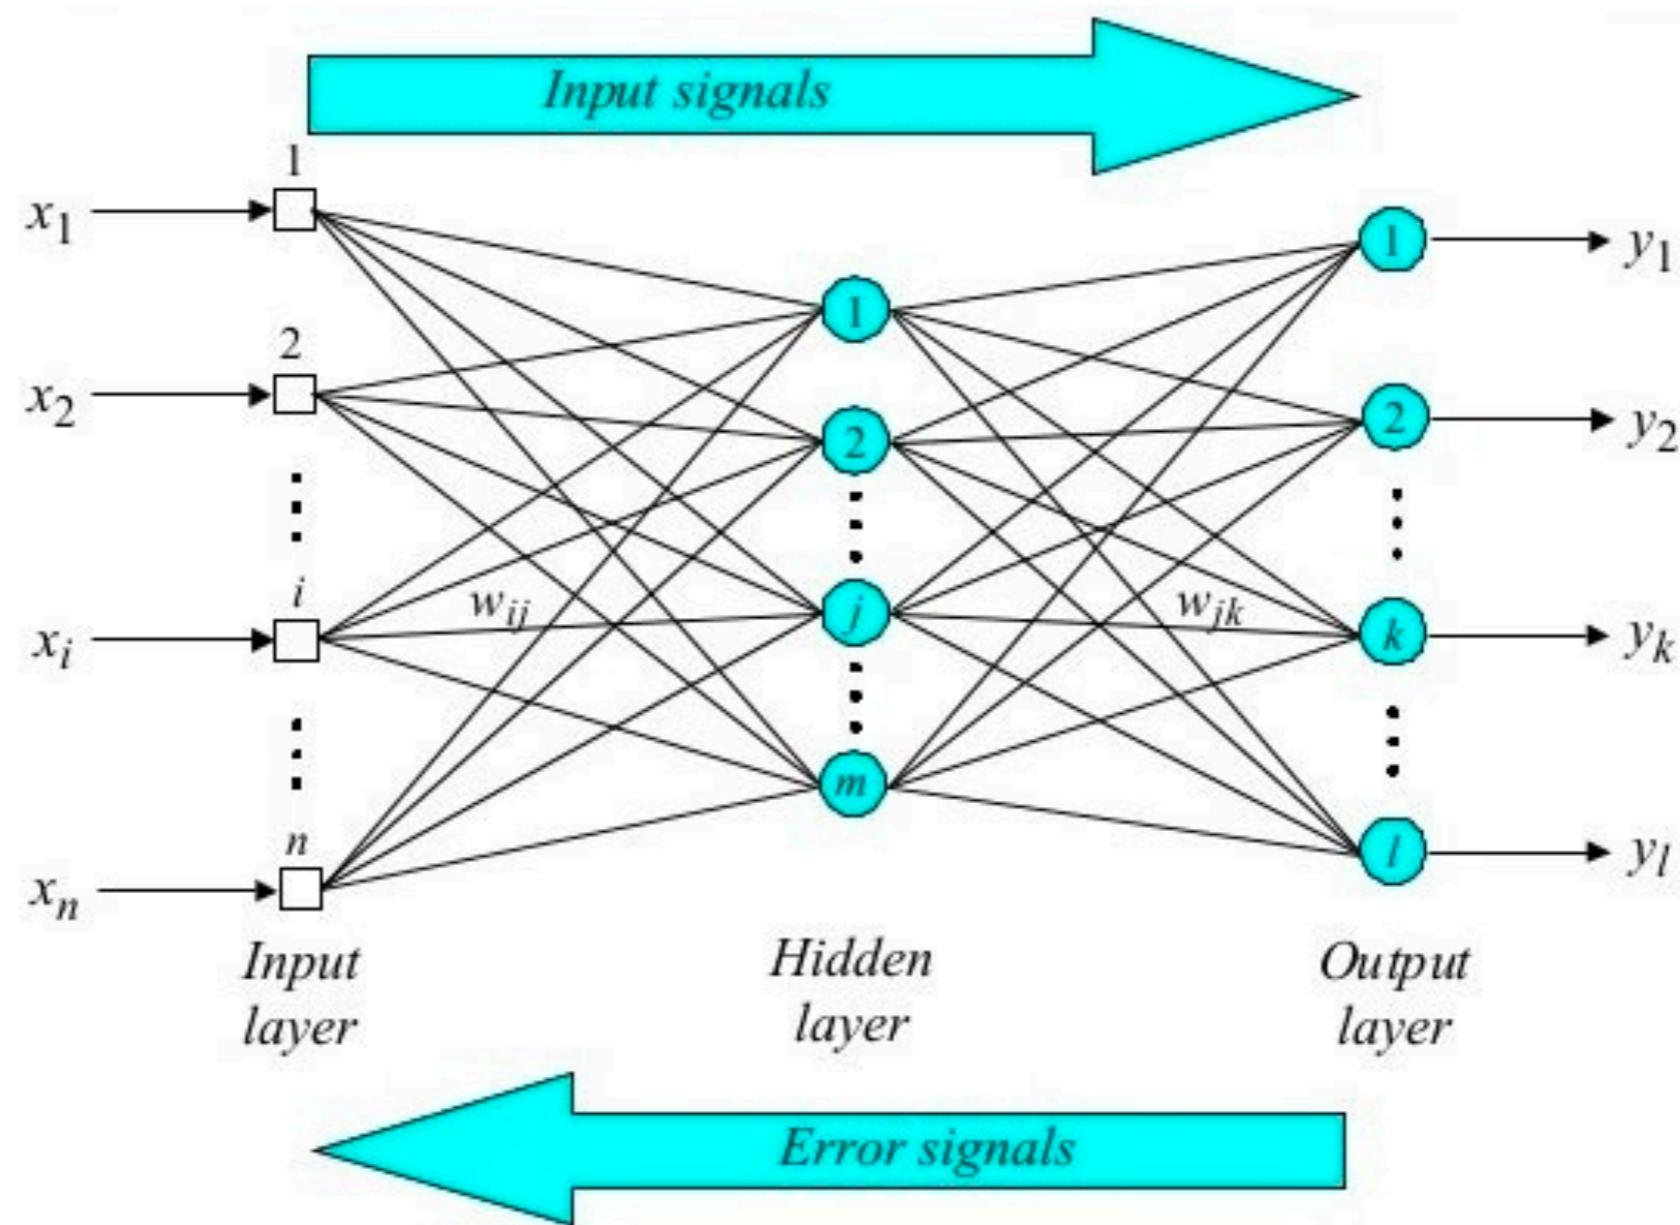
# Multilayer Neural Networks



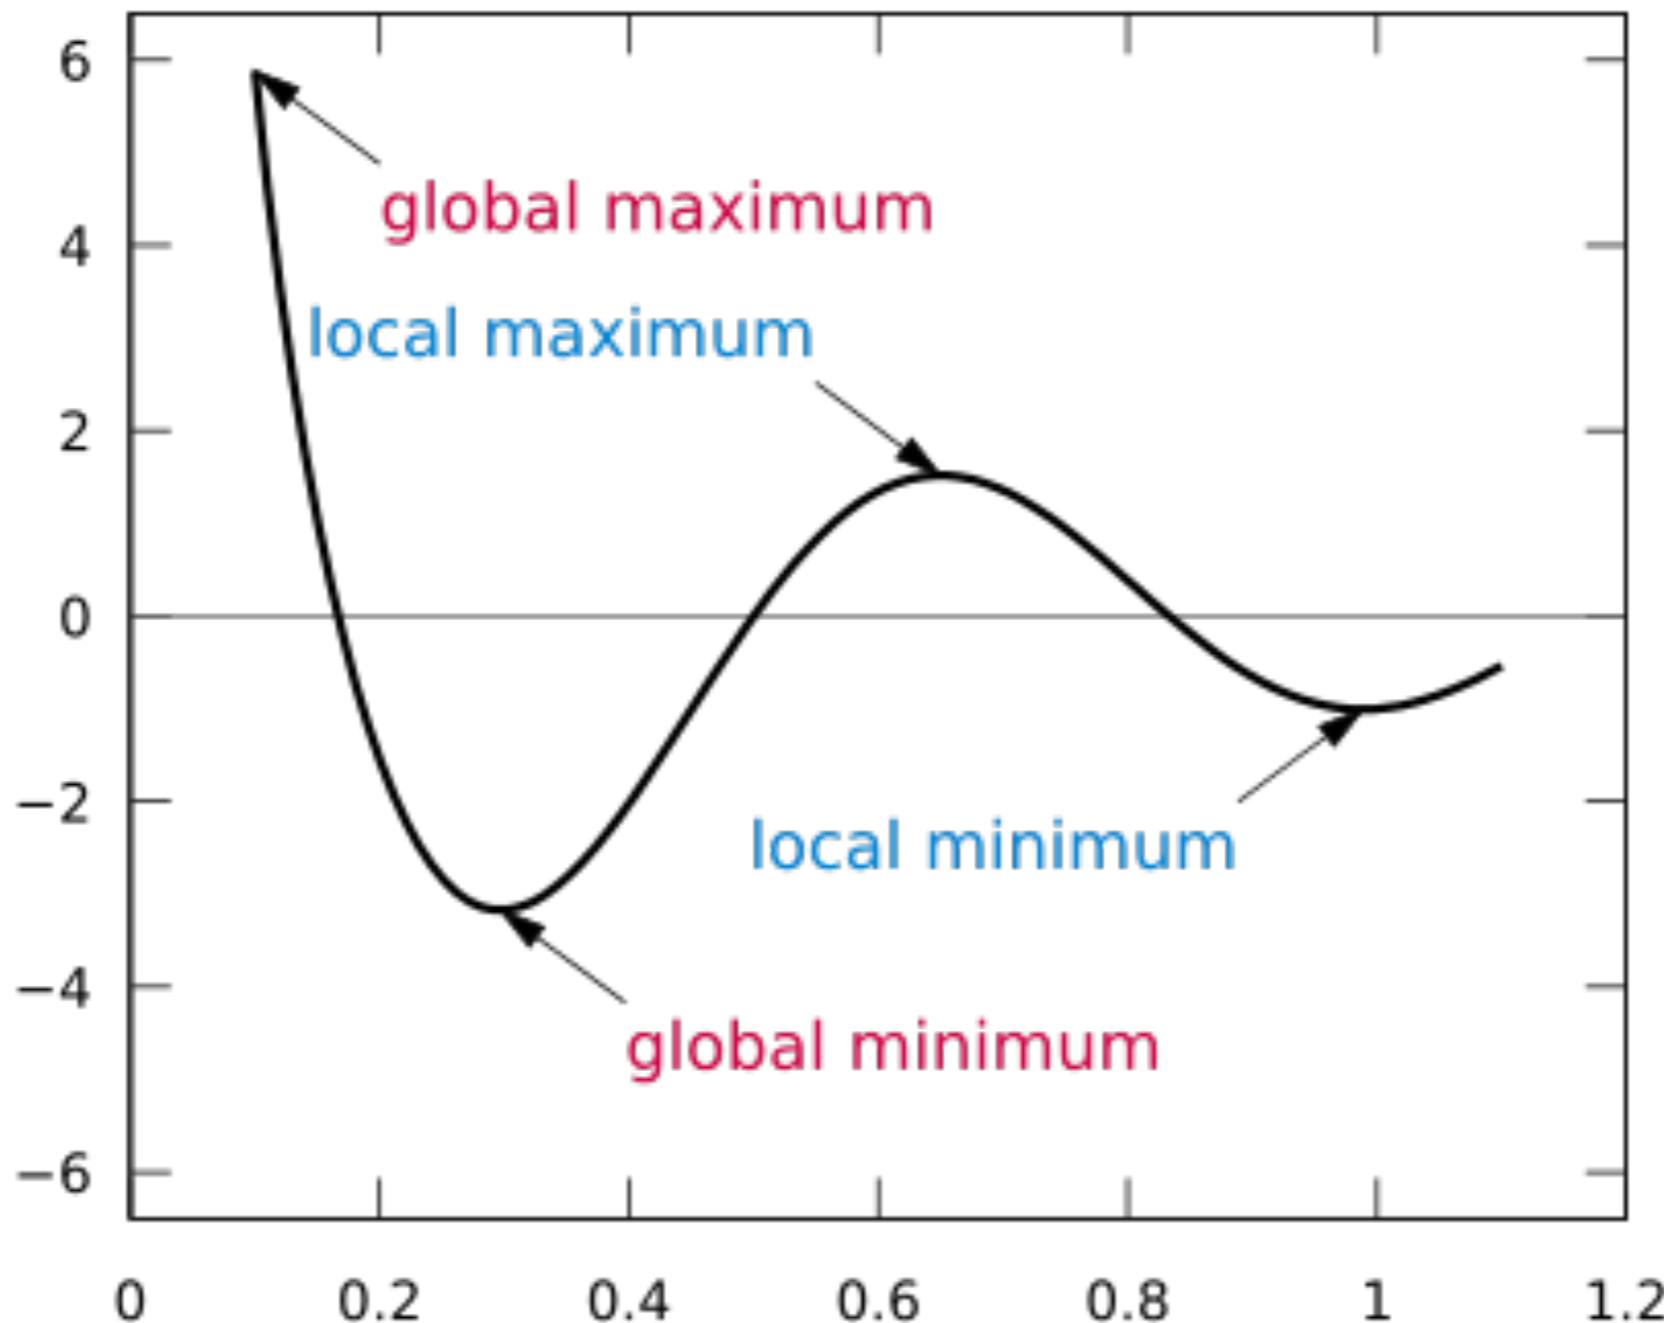
# Steepest gradient descent



# Back propagation

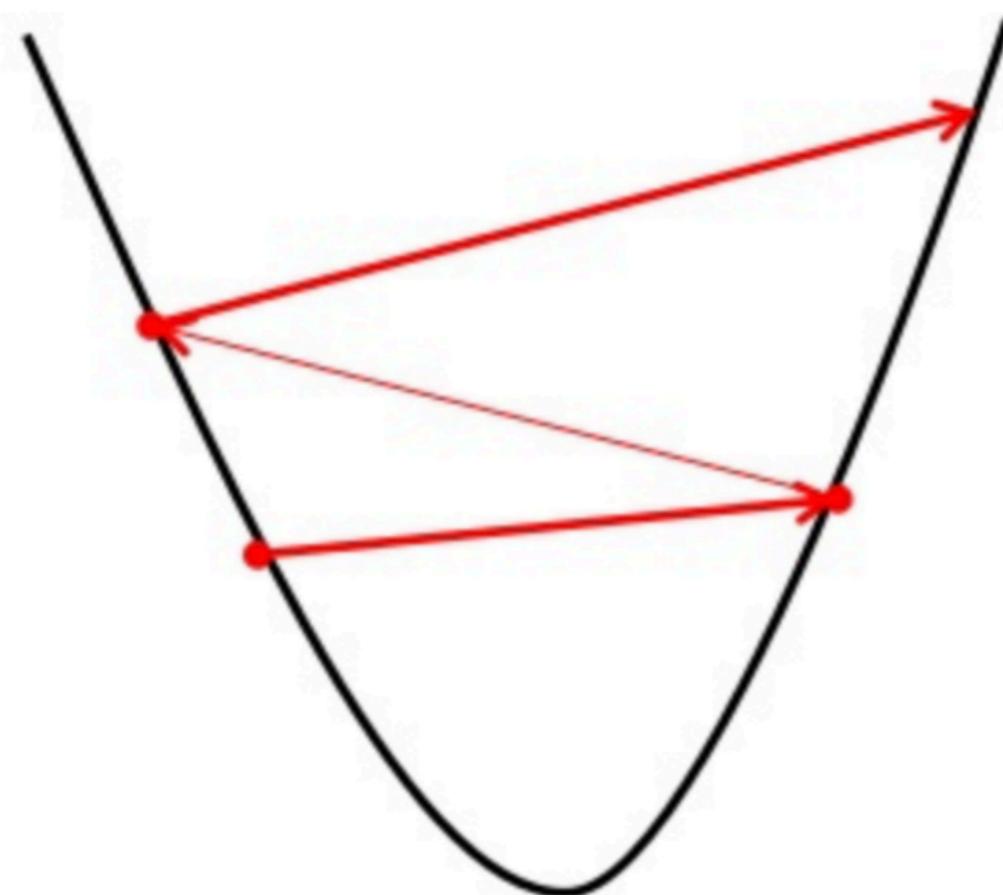


# Parameter optimization strategies

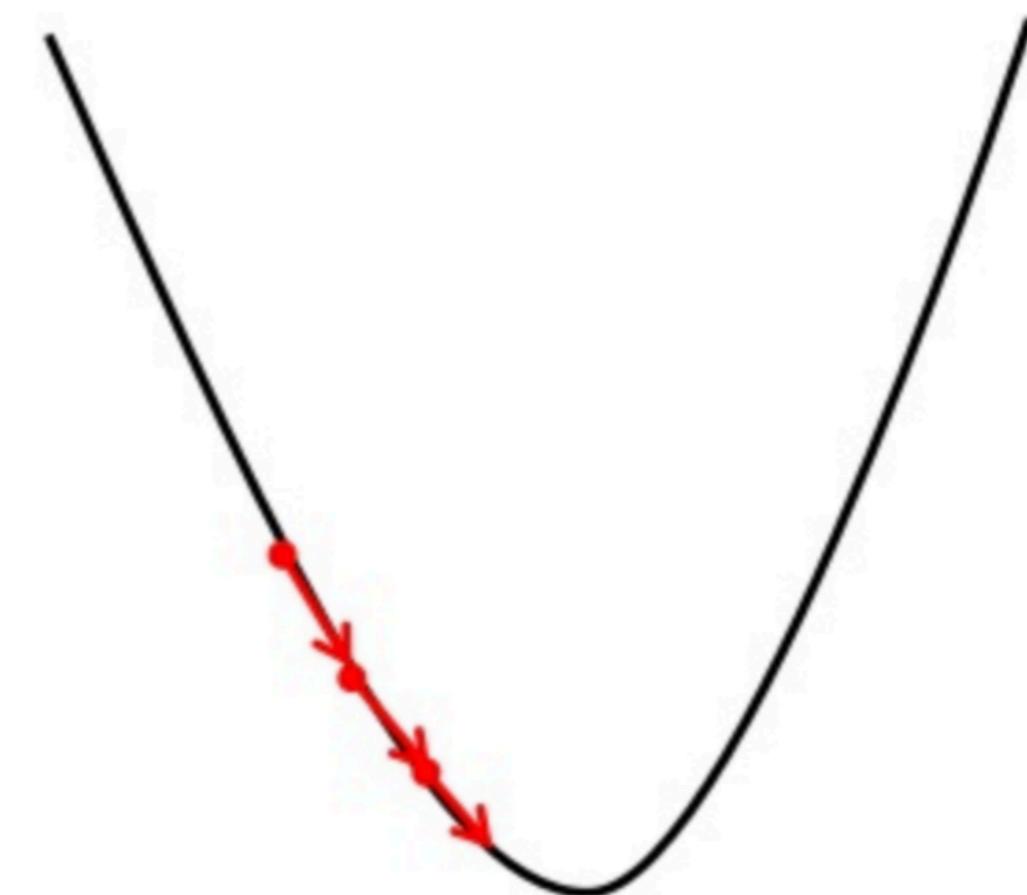


# Learning rate tuning

Big learning rate



Small learning rate



# Gradient Descent Variants

**(Batch) Gradient Descent**

$$w_{t+1} = w_t - \lambda \frac{\partial e(X, y)}{\partial w_t}$$

**Stochastic Gradient Descent**

$$w_{t+1} = w_t - \lambda \frac{\partial e(X^i, y^i)}{\partial w_t}$$

**Mini-Batch Gradient Descent**

$$w_{t+1} = w_t - \lambda \frac{\partial e(X^{(i,i+n)}, y^{(i,i+n)})}{\partial w_t}$$

# Momentum and Nesterov Accelerated Gradient

$$v_t = \gamma v_{t-1} + \lambda \frac{\partial e(w_t)}{\partial w_t}$$

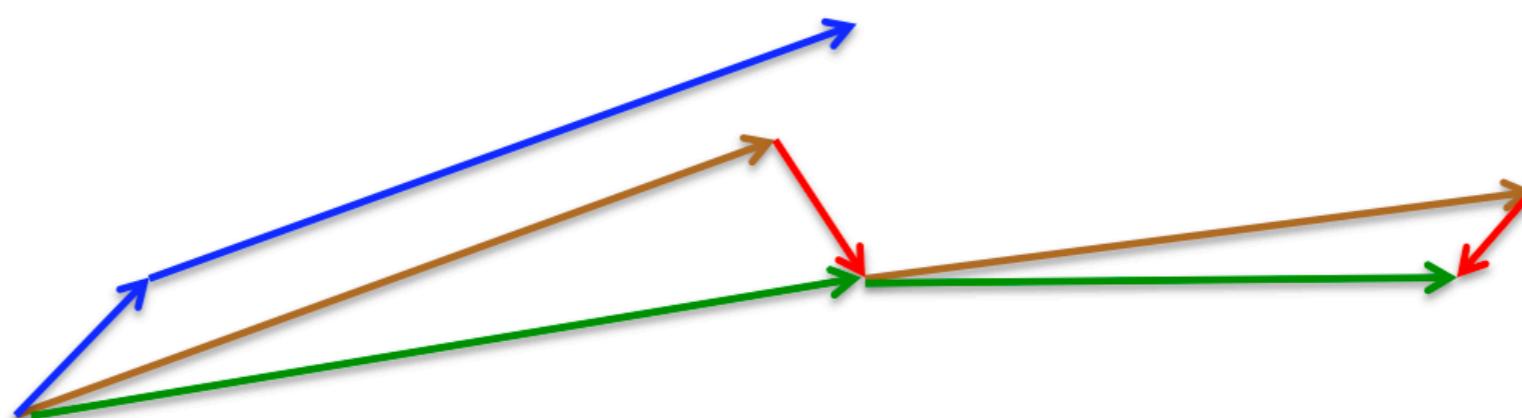
**Naive momentum**

$$w_{t+1} = w_t - v_t$$

**Nesterov Accelerated Gradient**

$$v_t = \gamma v_{t-1} + \lambda \frac{\partial e(w_t - \gamma v_{t-1})}{\partial w_t}$$

$$w_{t+1} = w_t - v_t$$



# Adaptive Gradient Algorithms

$$w_{t+1} = w_t - \frac{\lambda}{\sqrt{\sum_{i=1}^t g_i^2 + \epsilon}} g_t$$

**Adagrad**

$$g_i = \frac{\partial e(w_i)}{\partial w_i}$$

$$w_{t+1} = w_t - \frac{\lambda}{\sqrt{\mathbb{E}[g^2]_t - \epsilon}} g_t^2$$

**RMSProp**

$$\mathbb{E}[g^2]_t = \gamma \mathbb{E}[g^2]_{t-1} + (1 - \gamma) g_t^2$$

# Adam and Nadam

**Adam**

Combination of RMSProp with momentum

**Nadam**

Combination of RMSProp with Nesterov momentum

# Loss functions for deep learning

**Mean Squared Error**

$$MSE = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$

**Mean Absolute Error**

$$MSE = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$

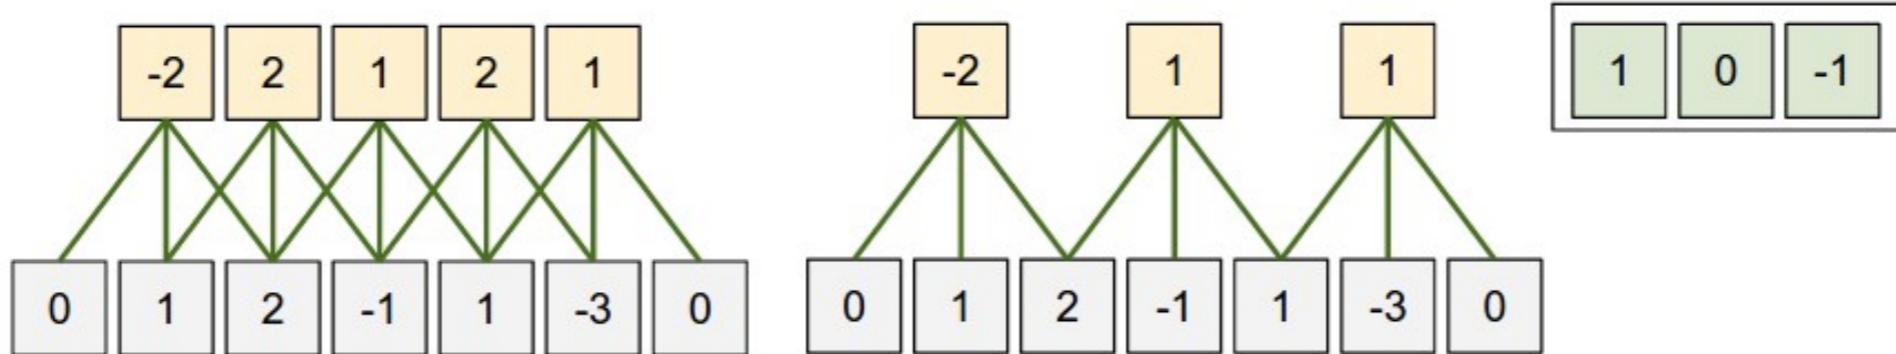
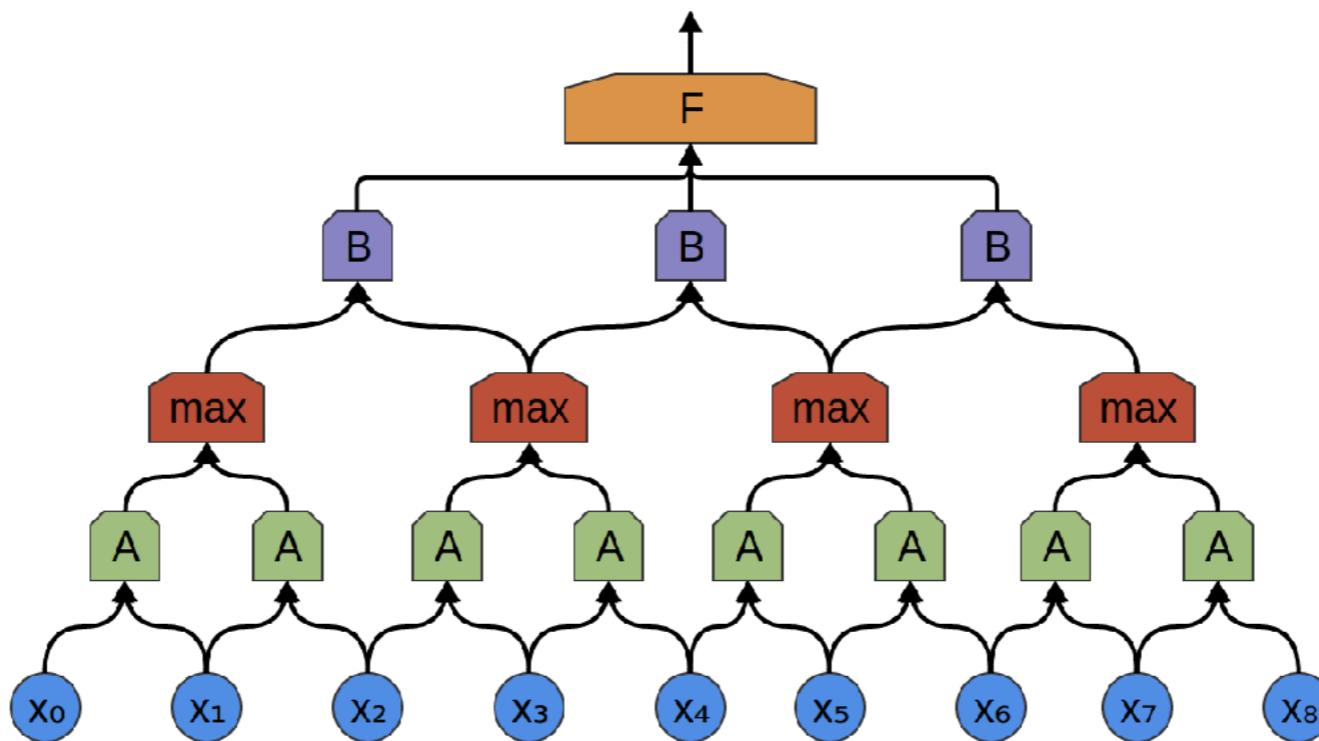
# Cross Entropy (Negative Log Likelihood)

**Categorical Cross Entropy** 
$$CCE = -\frac{\sum_{i=1}^n \sum_{j=1}^c y_{i,j} \log(\hat{y}_{i,j})}{n}$$

## Binary Cross Entropy

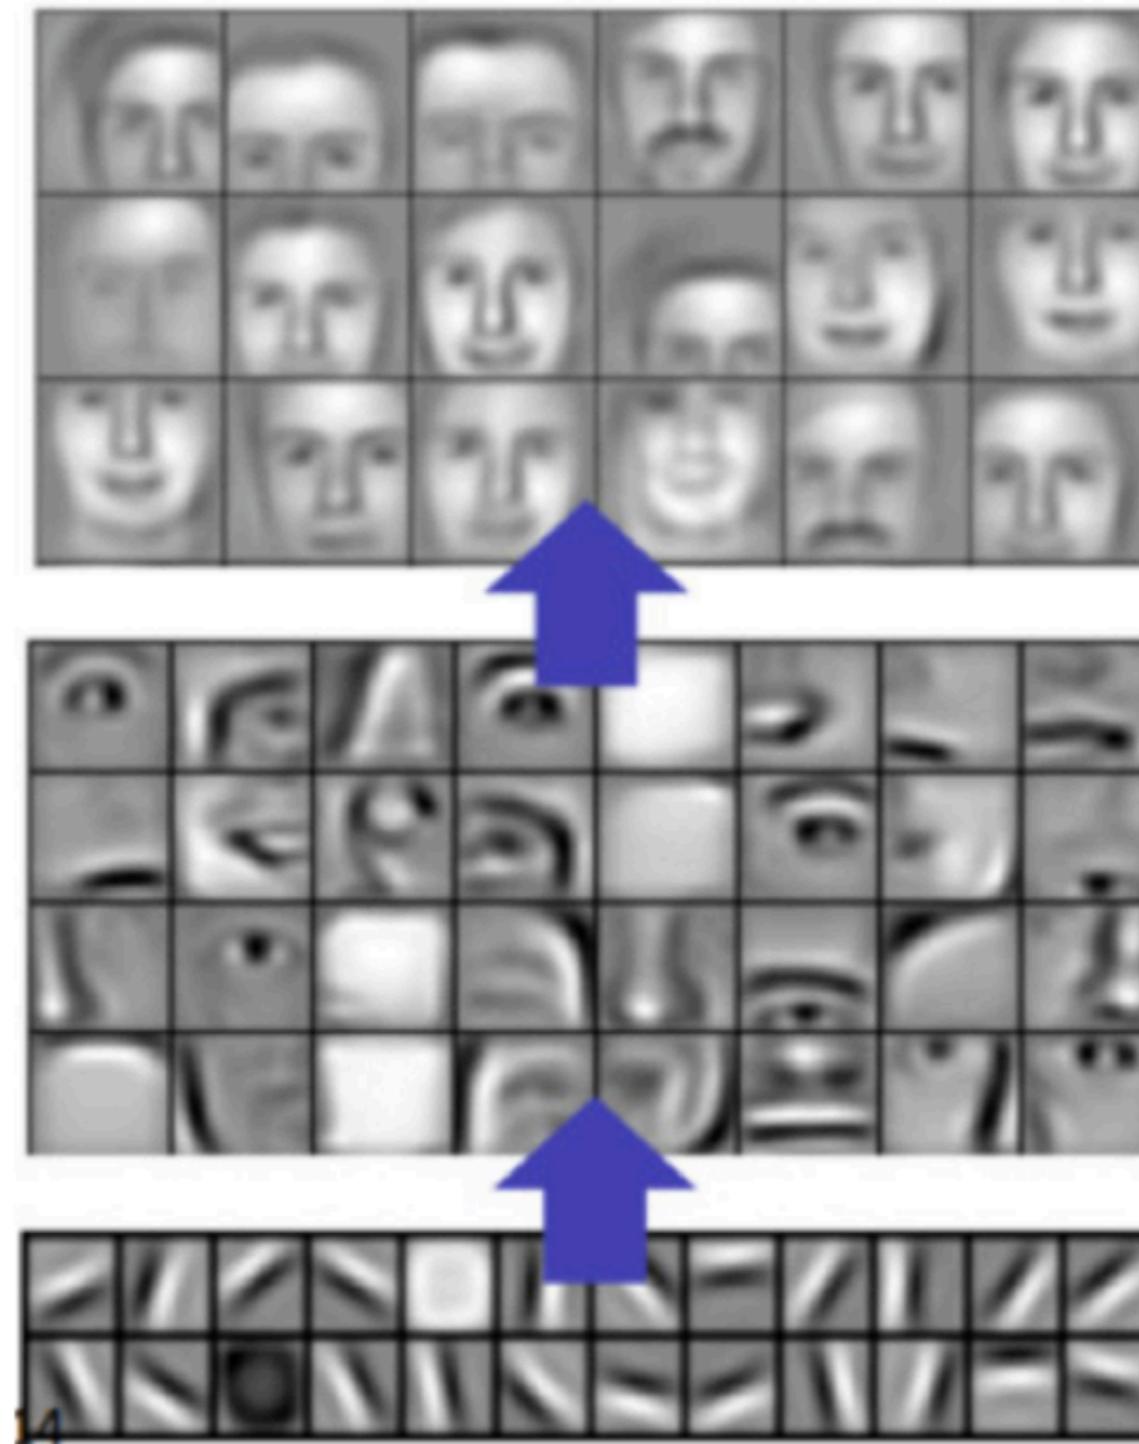
$$BCE = -\frac{\sum_{i=1}^n [y_{i,j} \log(\hat{y}_{i,j}) + (1 - y_{i,j}) \log(1 - \hat{y}_{i,j})]}{n}$$

# Convolution



Source: <https://www.tensorflow.org>

# Weights visualization



Layer 3

Layer 2

Layer 1

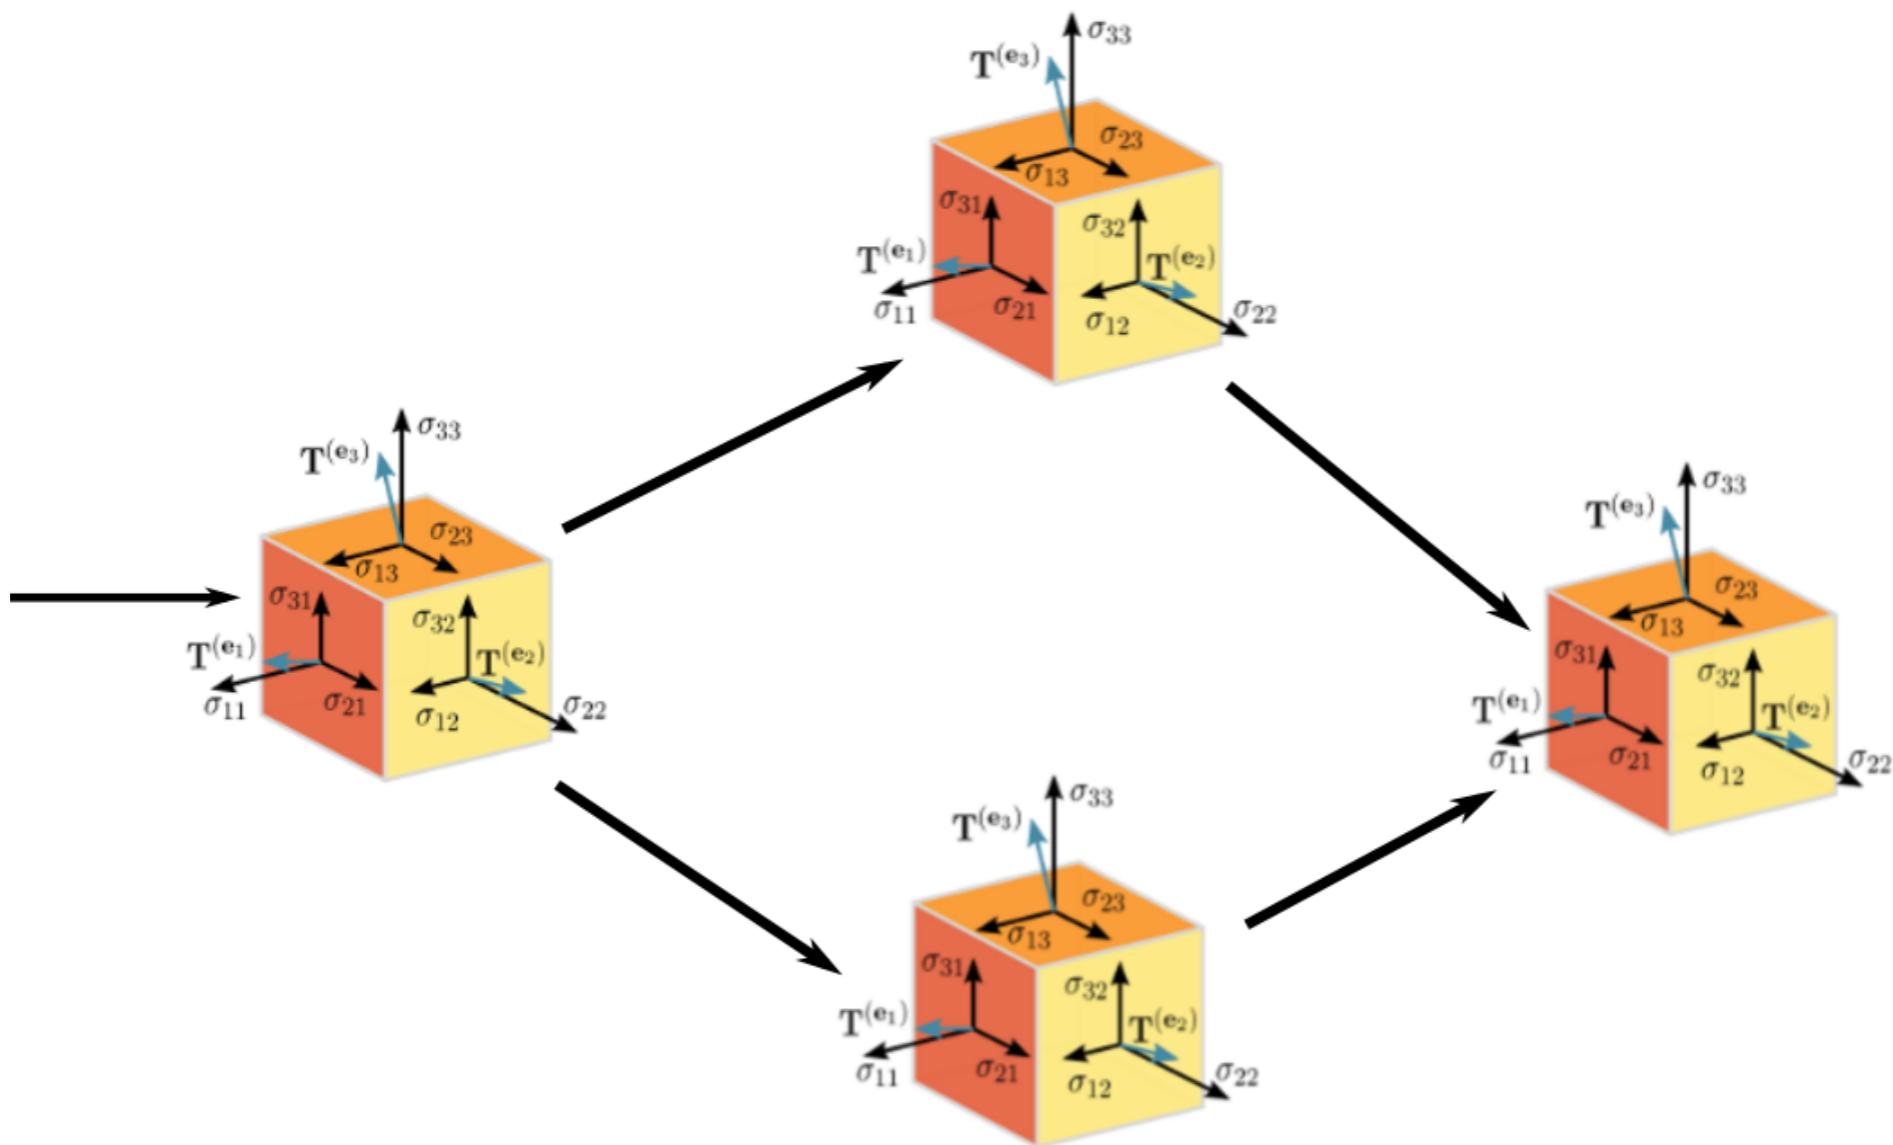
# Important terms

- deep learning
- stochastic gradient descent
- batch and mini-batch learning
- epoch
- dropout

# What is not TensorFlow



# What is TensorFlow?



# Keras tutorial

[\*\*Keras-introduction.ipynb\*\*](#)

# Neural Network architectures design



# Neural Network design best practices

- ★ Start from simple architectures
- ★ Get inspiration from architectures for similar problems
- ★ Change one parameter only and then validate

# Common architectures

**Feed forward network**

**Convolutional network**

**Recurrent network**

**Autoencoder network and Restricted Boltzmann Machines**

**Transformer network**

**U-Net**

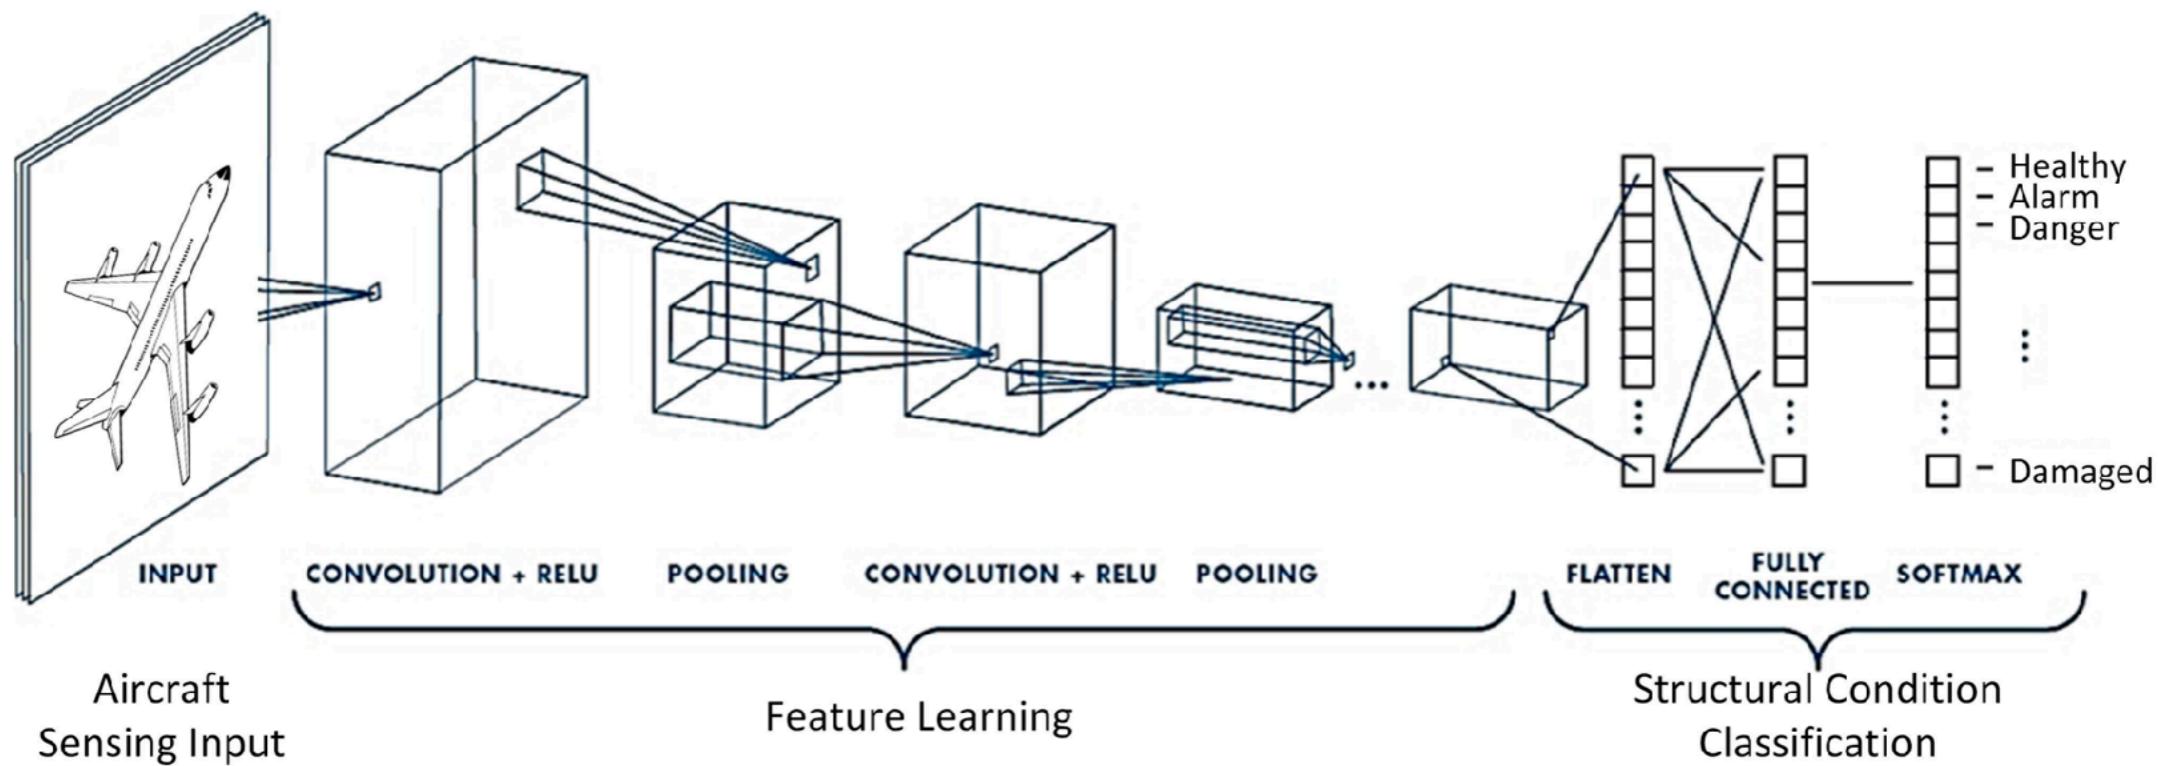
**Generative adversarial network**

# Implementation of a classification and regression task using NN

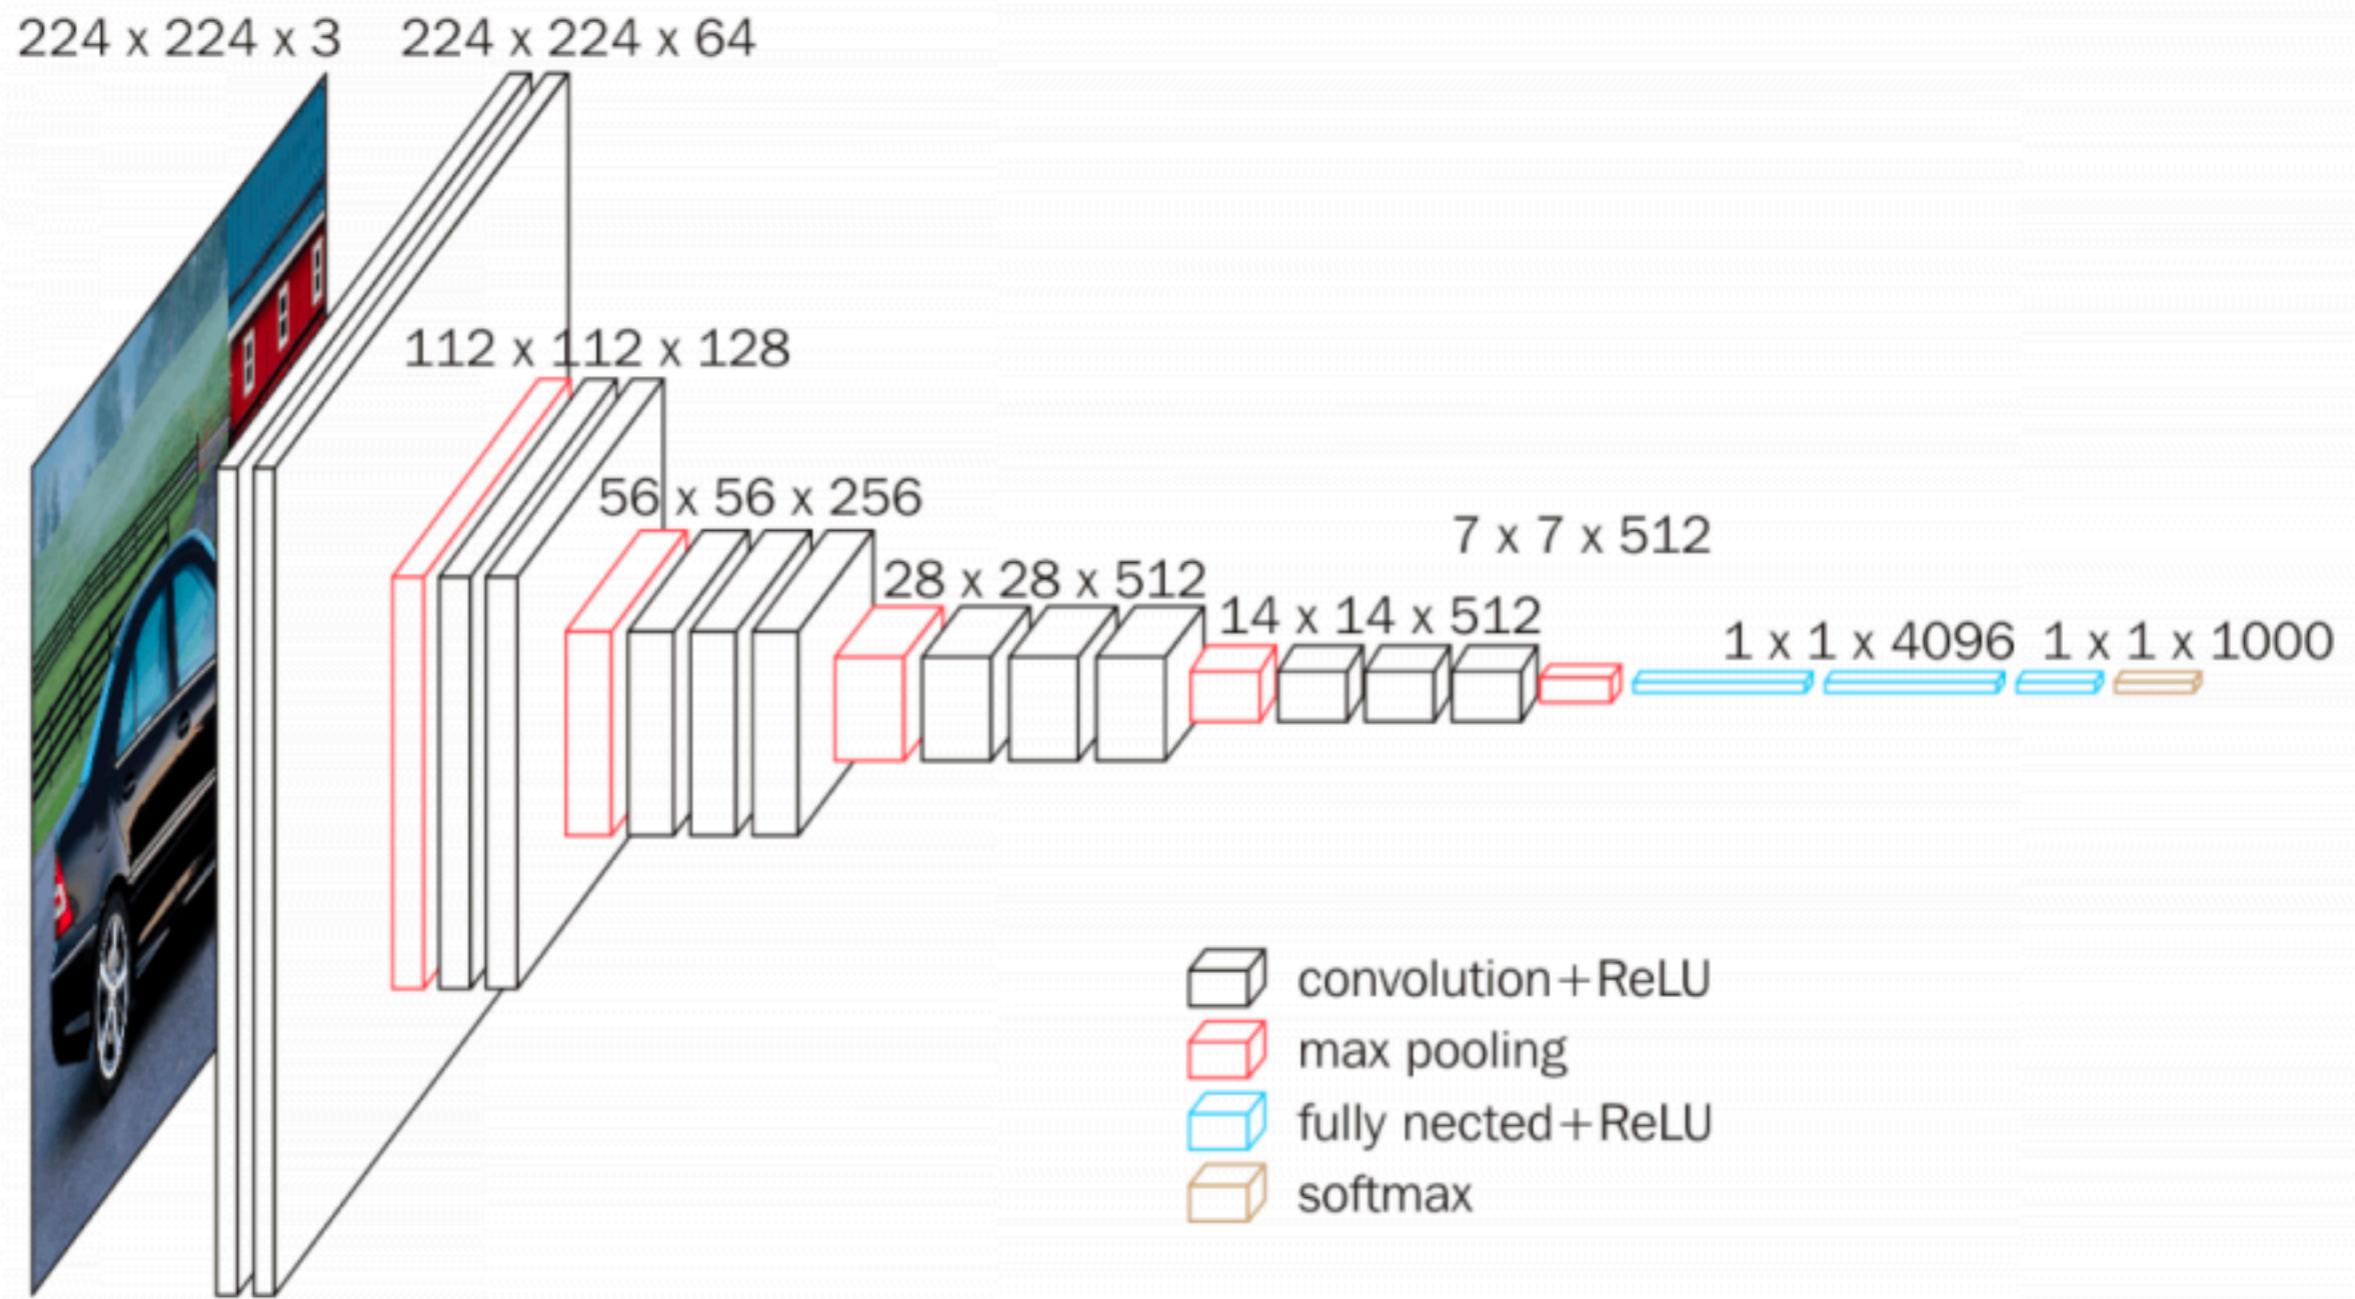
[\*\*Classification-nn-assignment.ipynb\*\*](#)

[\*\*Regression-nn-assignment.ipynb\*\*](#)

# Convolutional Neural Network



# VGG 16

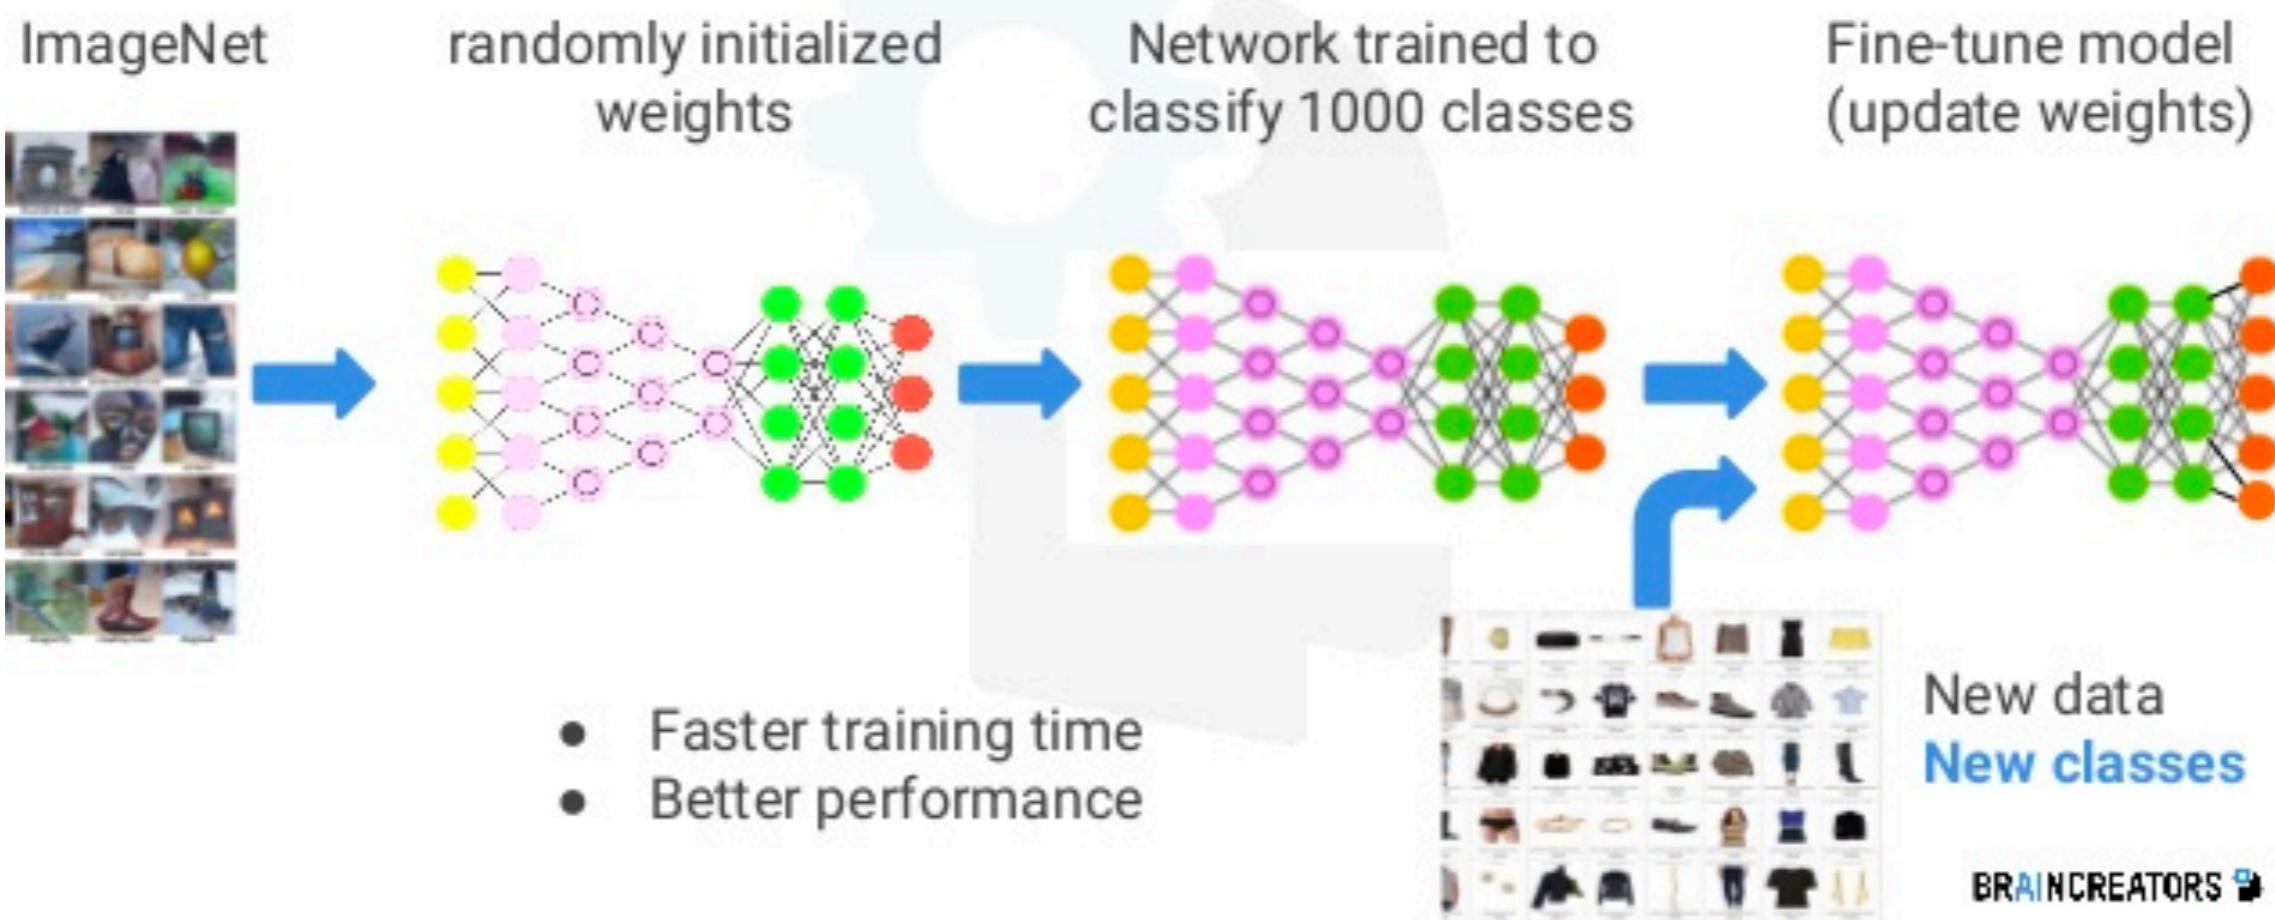


# ResNet



# Finetuning

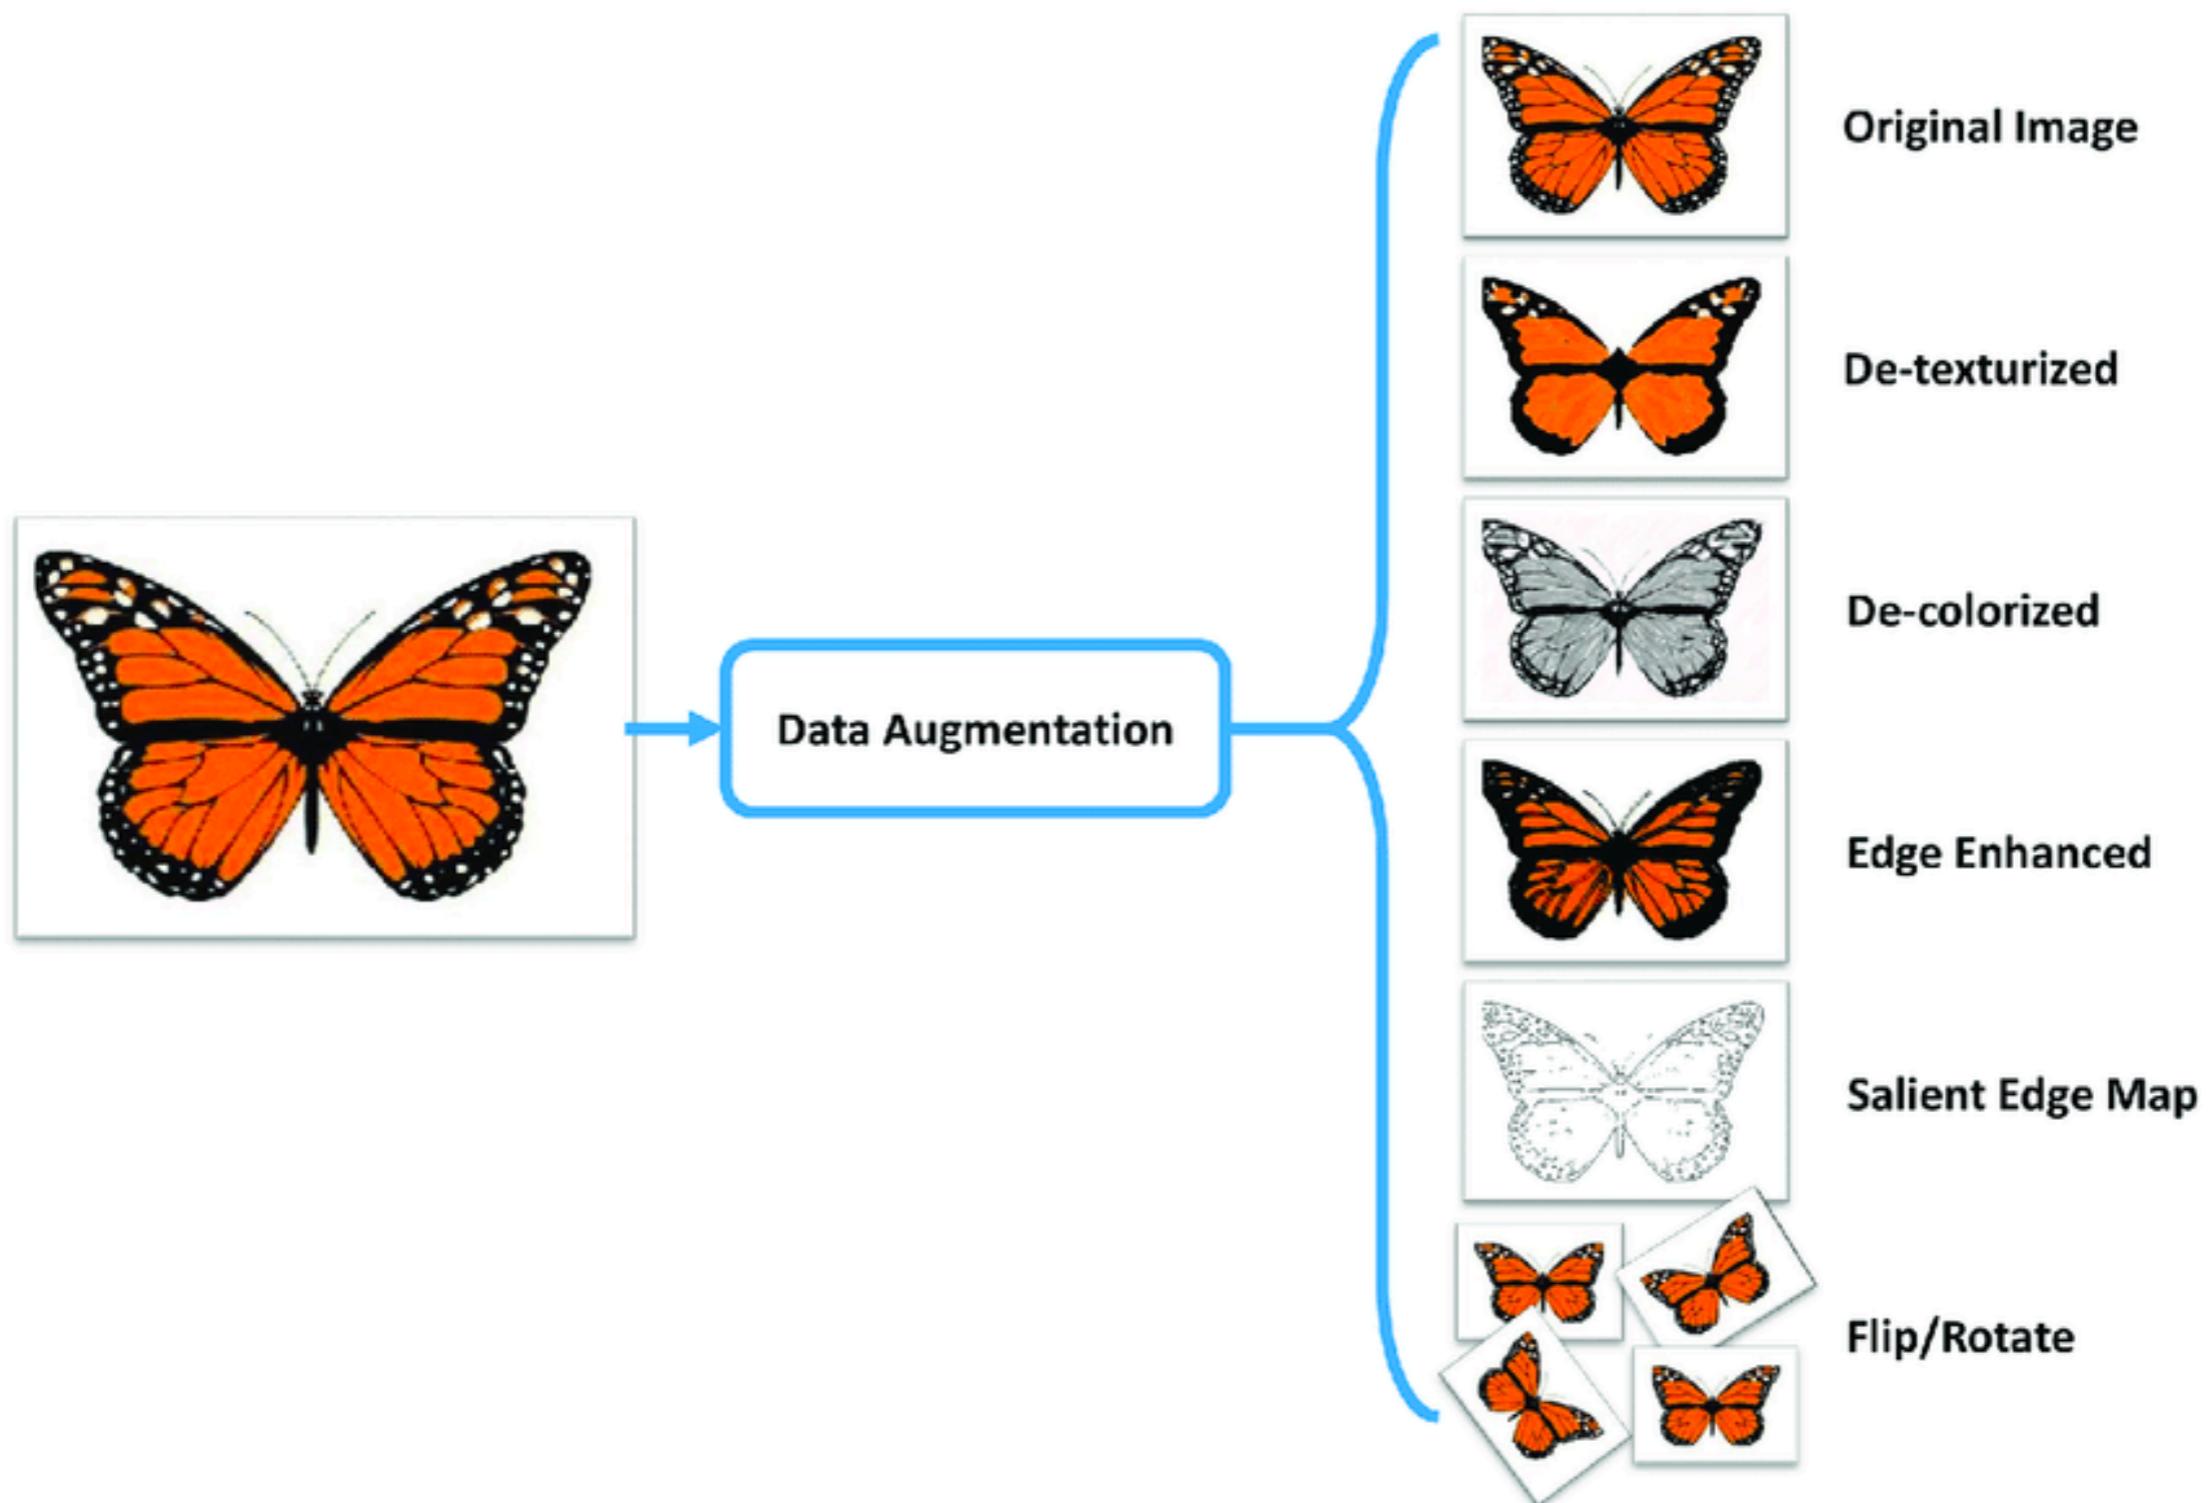
## Transfer Learning



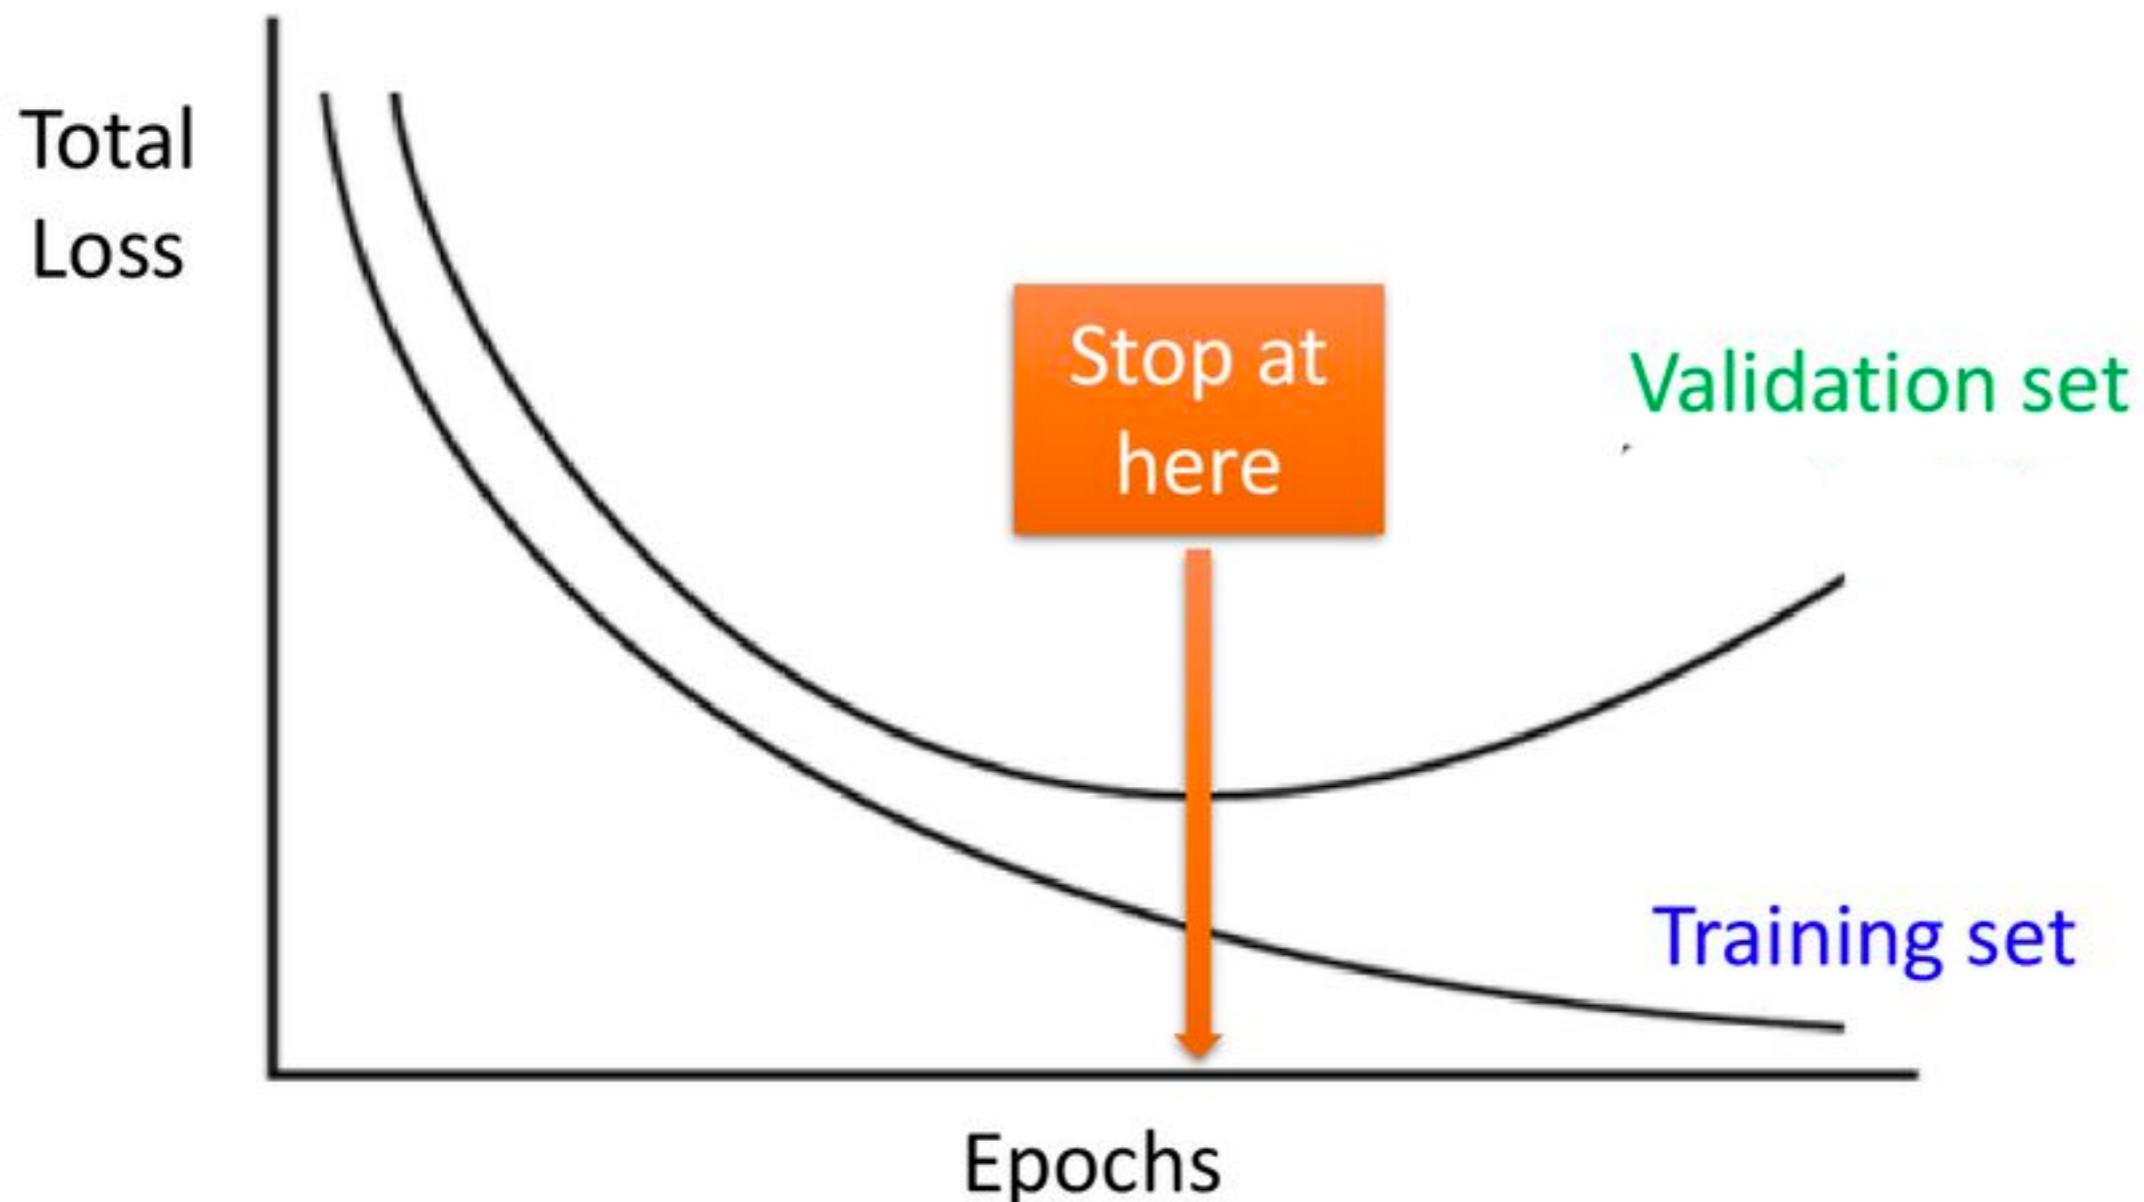
# **Transfer learning example**

**Transfer\_learning.ipynb**

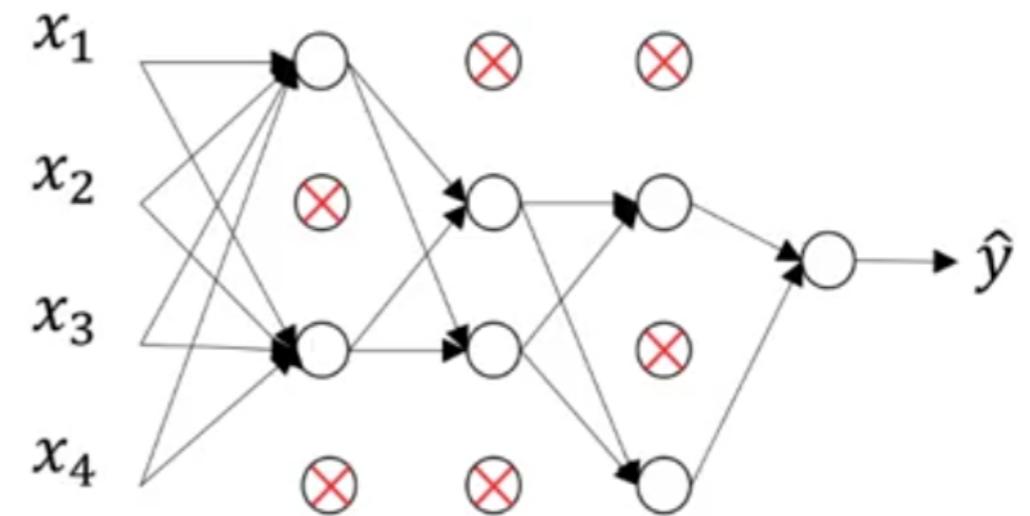
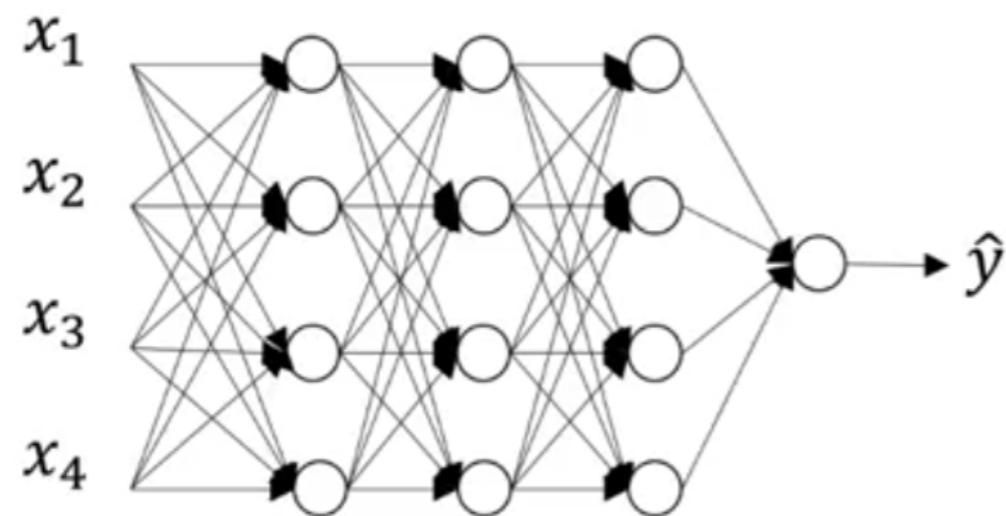
# Data augmentation



# Early stopping



# Dropout



# L2 Regularization in deep learning

$$cost(w^1, b^1, \dots, w^L, b^L) = \frac{1}{n} \sum_{i=1}^n Loss(y_i, \hat{y}_i) + \frac{\lambda}{2n} \sum_{l=1}^L \|w^l\|_F^2$$

**Frobenius norm**

$$\|w\|_F = \sqrt{\sum_{i=1}^n \sum_{j=1}^m |w_{i,j}|^2}$$

# Batch normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

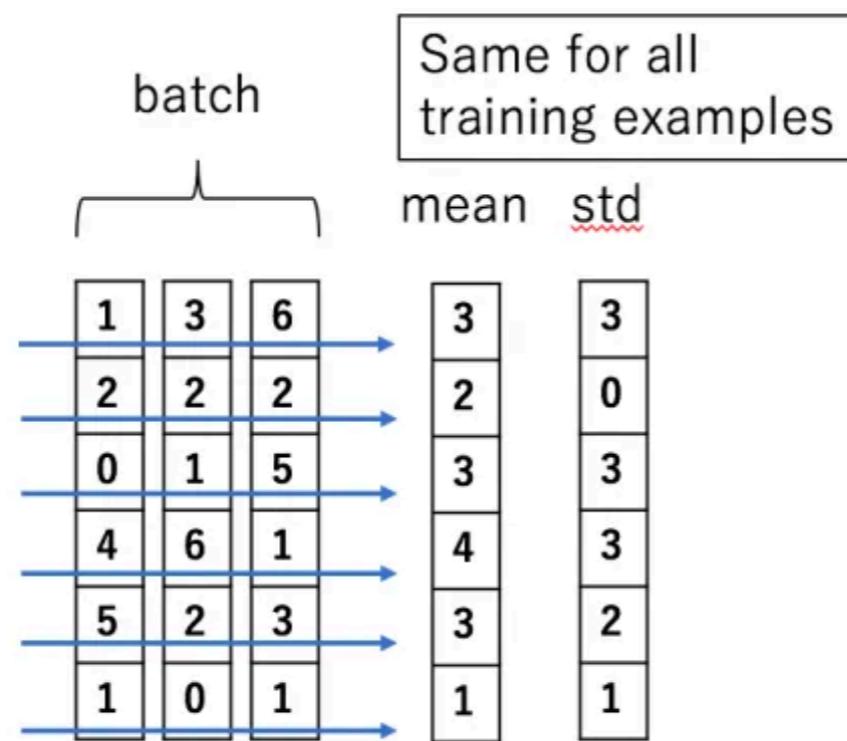
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

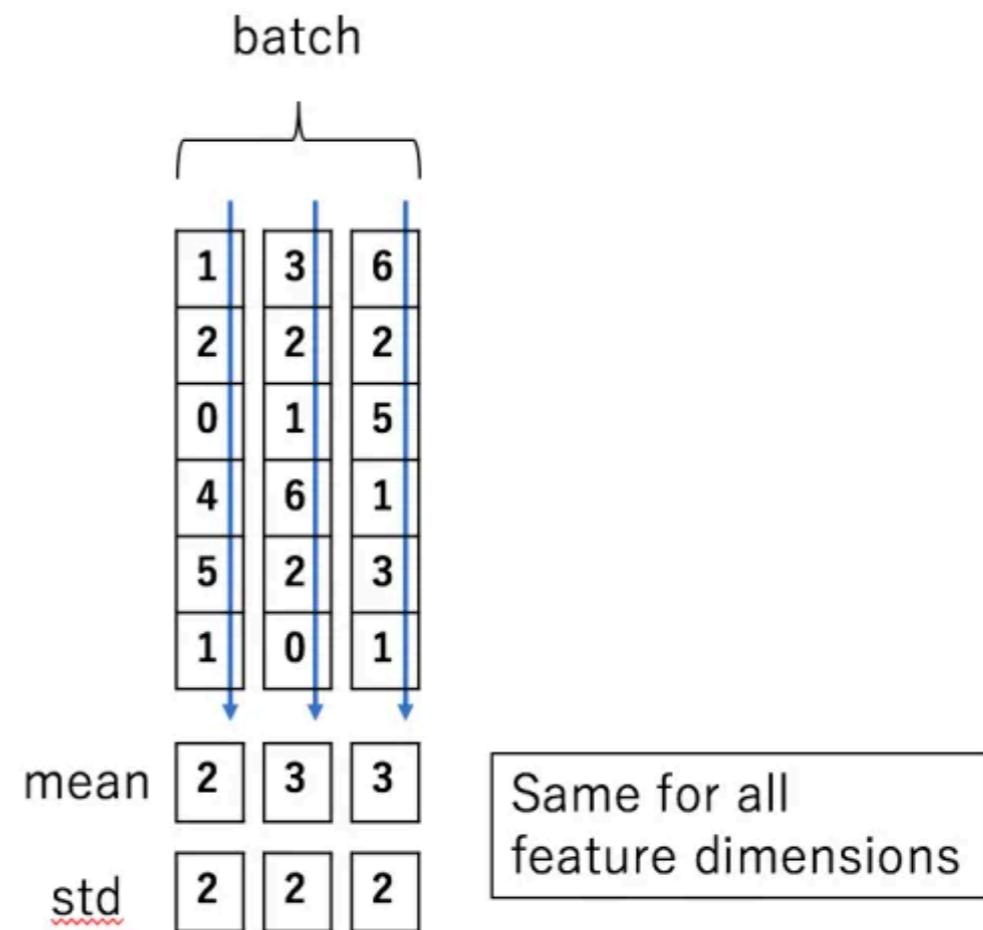
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

# Layer normalization

Batch Normalization



Layer Normalization



# Functional API in Keras

```
1 # Sequential model
2 from keras.models import Sequential
3 from keras.layers import Dense
4
5 model = Sequential()
6 model.add(10, input_shape=(10,), activation='relu')
7 model.add(Dense(20, activation='relu'))
8 model.add(Dense(10, activation='relu'))
9 model.add(Dense(1, activation='sigmoid'))
```

```
1 # Functional model
2 from keras.models import Model
3 from keras.layers import Input, Dense
4
5 visible = Input(shape=(10,))
6 hidden1 = Dense(10, activation='relu')(visible)
7 hidden2 = Dense(20, activation='relu')(hidden1)
8 hidden3 = Dense(10, activation='relu')(hidden2)
9 output = Dense(1, activation='sigmoid')(hidden3)
10 model = Model(inputs=visible, outputs=output)
```

# Shared Input

```
1 # Shared Input Layer
2 from keras.utils import plot_model
3 from keras.models import Model
4 from keras.layers import Input, Dense, Flatten
5 from keras.layers.convolutional import Conv2D
6 from keras.layers.pooling import MaxPooling2D
7 from keras.layers.merge import concatenate
8 # input layer
9 visible = Input(shape=(64,64,1))
10 # first feature extractor
11 conv1 = Conv2D(32, kernel_size=4, activation='relu')(visible)
12 pool1 = MaxPooling2D(pool_size=(2, 2))(conv1)
13 flat1 = Flatten()(pool1)
14 # second feature extractor
15 conv2 = Conv2D(16, kernel_size=8, activation='relu')(visible)
16 pool2 = MaxPooling2D(pool_size=(2, 2))(conv2)
17 flat2 = Flatten()(pool2)
18 # merge feature extractors
19 merge = concatenate([flat1, flat2])
20 # interpretation layer
21 hidden1 = Dense(10, activation='relu')(merge)
22 # prediction output
23 output = Dense(1, activation='sigmoid')(hidden1)
24 model = Model(inputs=visible, outputs=output)
```

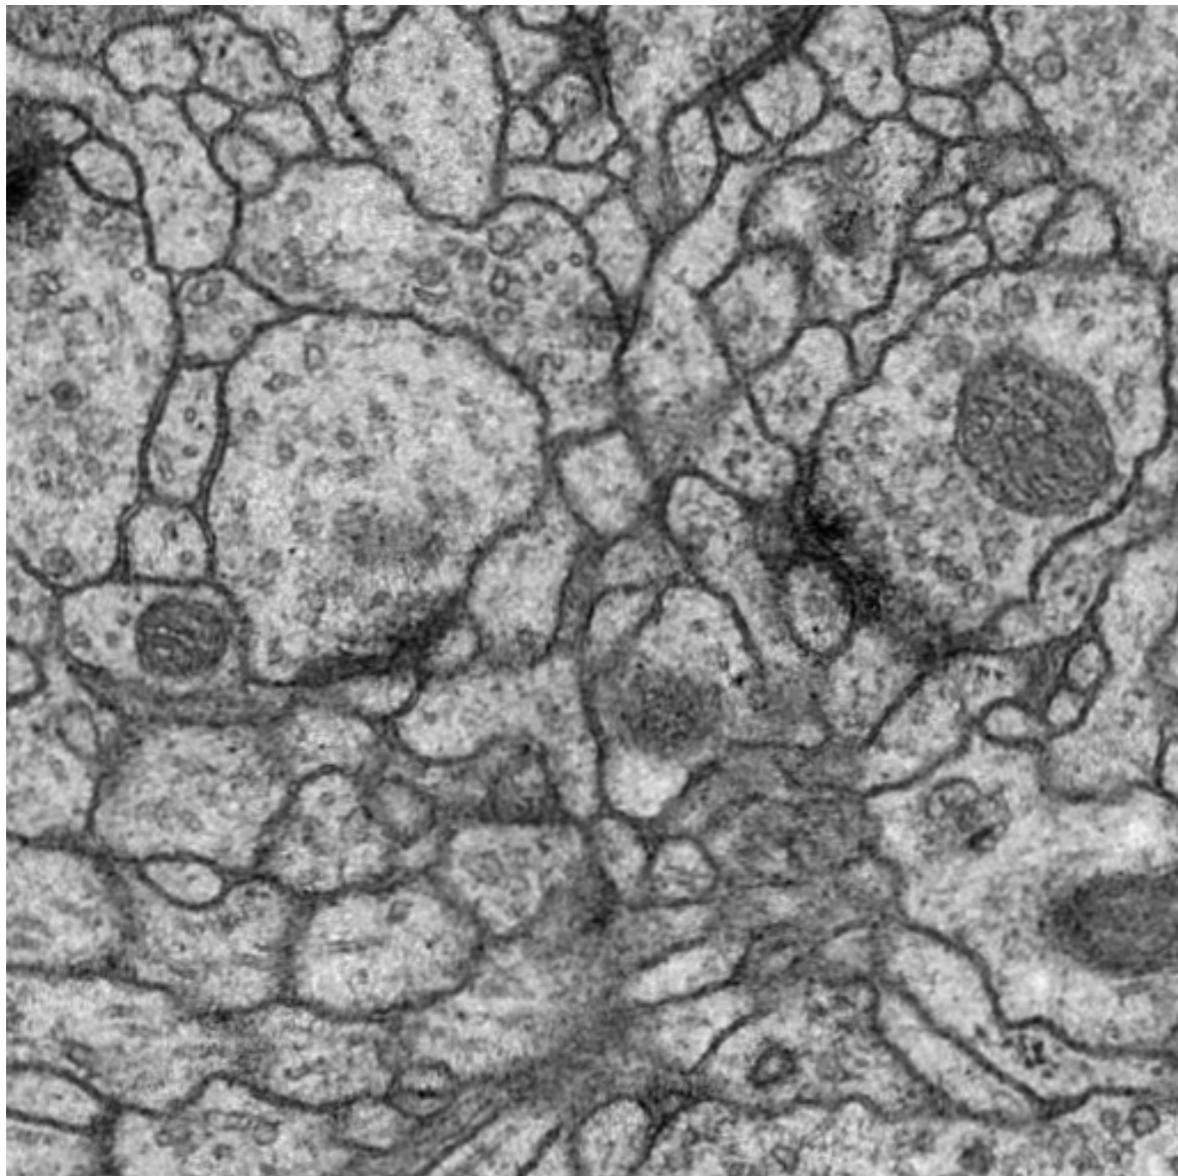
# Multiple inputs (outputs)

```
1 # Multiple Inputs
2 from keras.utils import plot_model
3 from keras.models import Model
4 from keras.layers import Input
5 from keras.layers import Dense
6 from keras.layers import Flatten
7 from keras.layers.convolutional import Conv2D
8 from keras.layers.pooling import MaxPooling2D
9 from keras.layers.merge import concatenate
10 # first input model
11 visible1 = Input(shape=(64,64,1))
12 conv11 = Conv2D(32, kernel_size=4, activation='relu')(visible1)
13 pool11 = MaxPooling2D(pool_size=(2, 2))(conv11)
14 conv12 = Conv2D(16, kernel_size=4, activation='relu')(pool11)
15 pool12 = MaxPooling2D(pool_size=(2, 2))(conv12)
16 flat1 = Flatten()(pool12)
17 # second input model
18 visible2 = Input(shape=(32,32,3))
19 conv21 = Conv2D(32, kernel_size=4, activation='relu')(visible2)
20 pool21 = MaxPooling2D(pool_size=(2, 2))(conv21)
21 conv22 = Conv2D(16, kernel_size=4, activation='relu')(pool21)
22 pool22 = MaxPooling2D(pool_size=(2, 2))(conv22)
23 flat2 = Flatten()(pool22)
24 # merge input models
25 merge = concatenate([flat1, flat2])
26 # interpretation model
27 hidden1 = Dense(10, activation='relu')(merge)
28 hidden2 = Dense(10, activation='relu')(hidden1)
29 output = Dense(1, activation='sigmoid')(hidden2)
30 model = Model(inputs=[visible1, visible2], outputs=output)
```

# Practical example on regularization and normalization

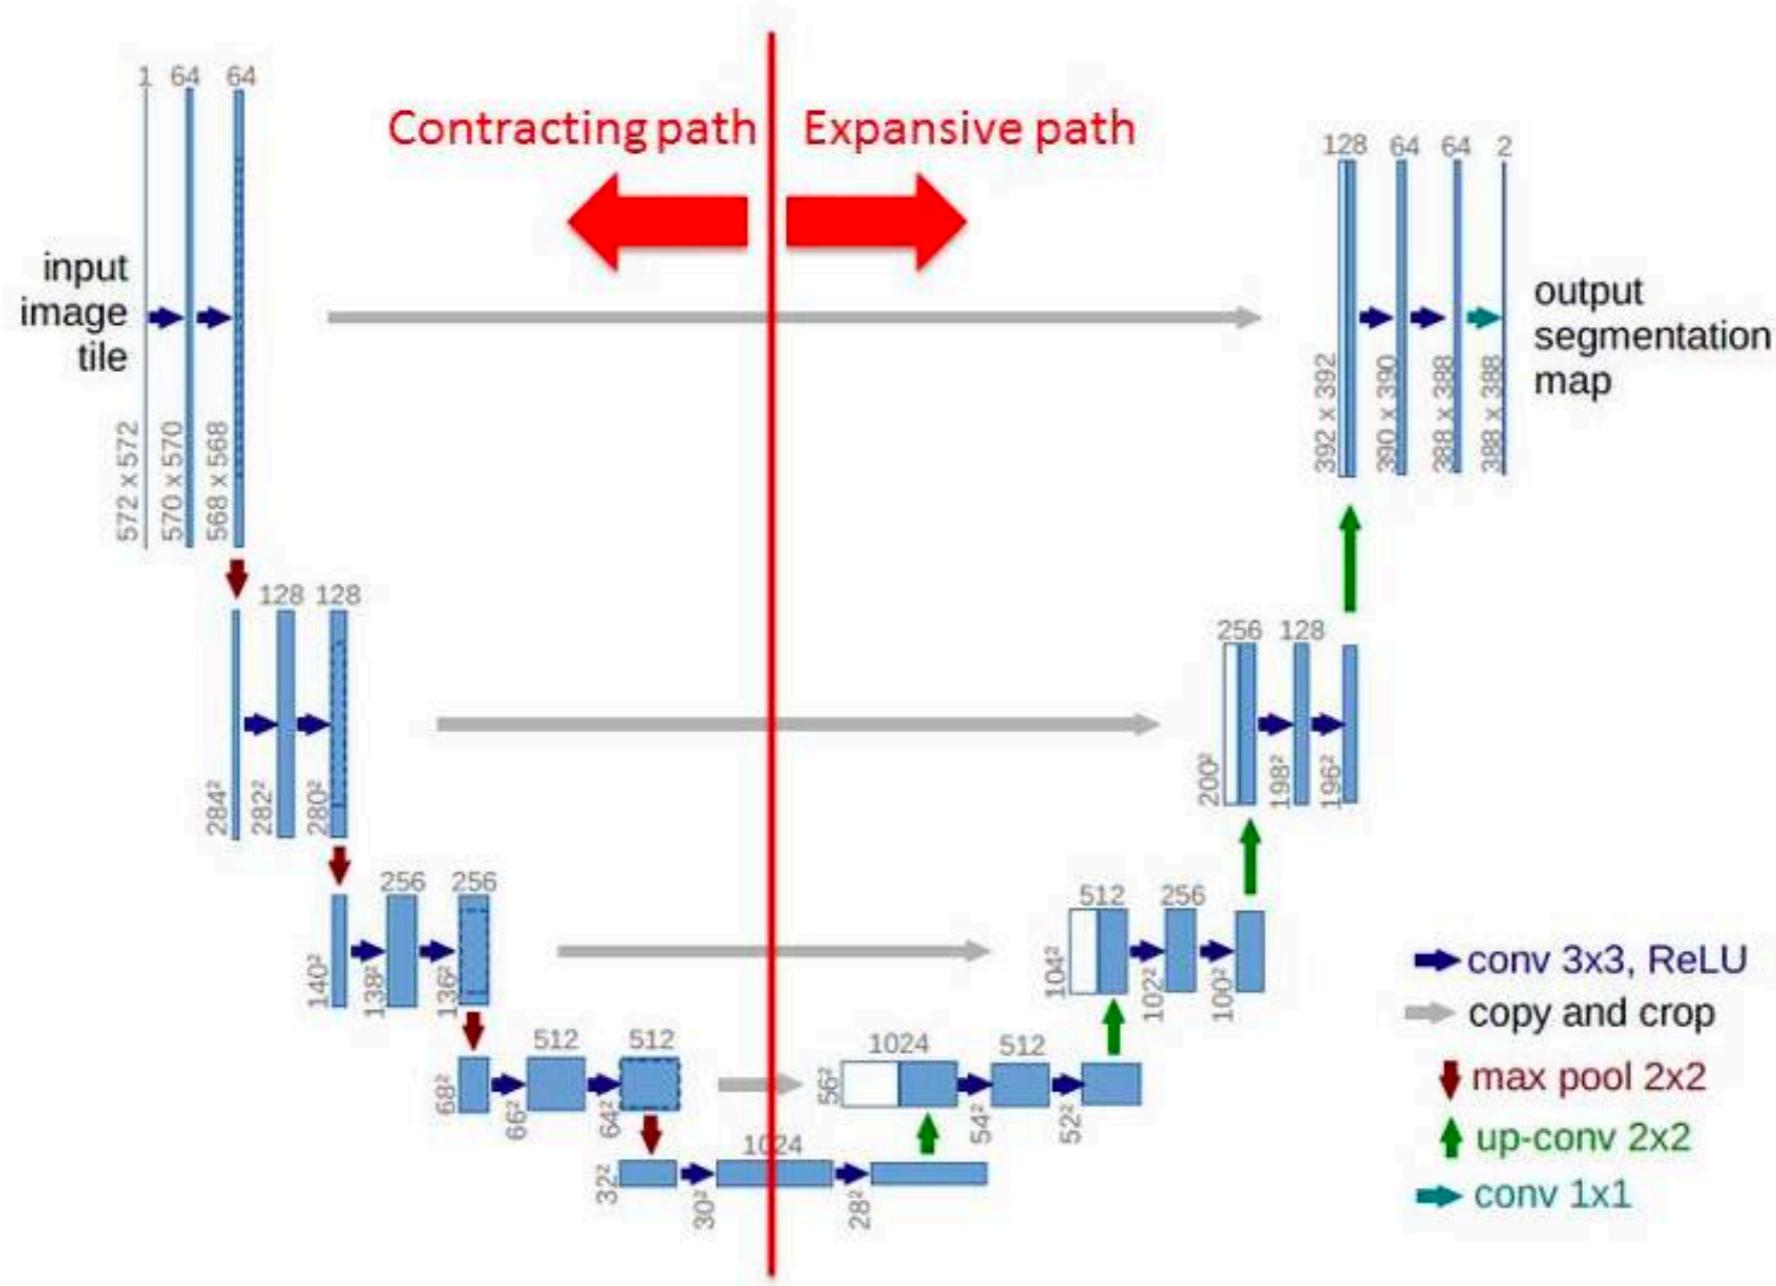
[Normalization-and-regularization-assignment.ipynb](#)

# Image segmentation



# U-Net

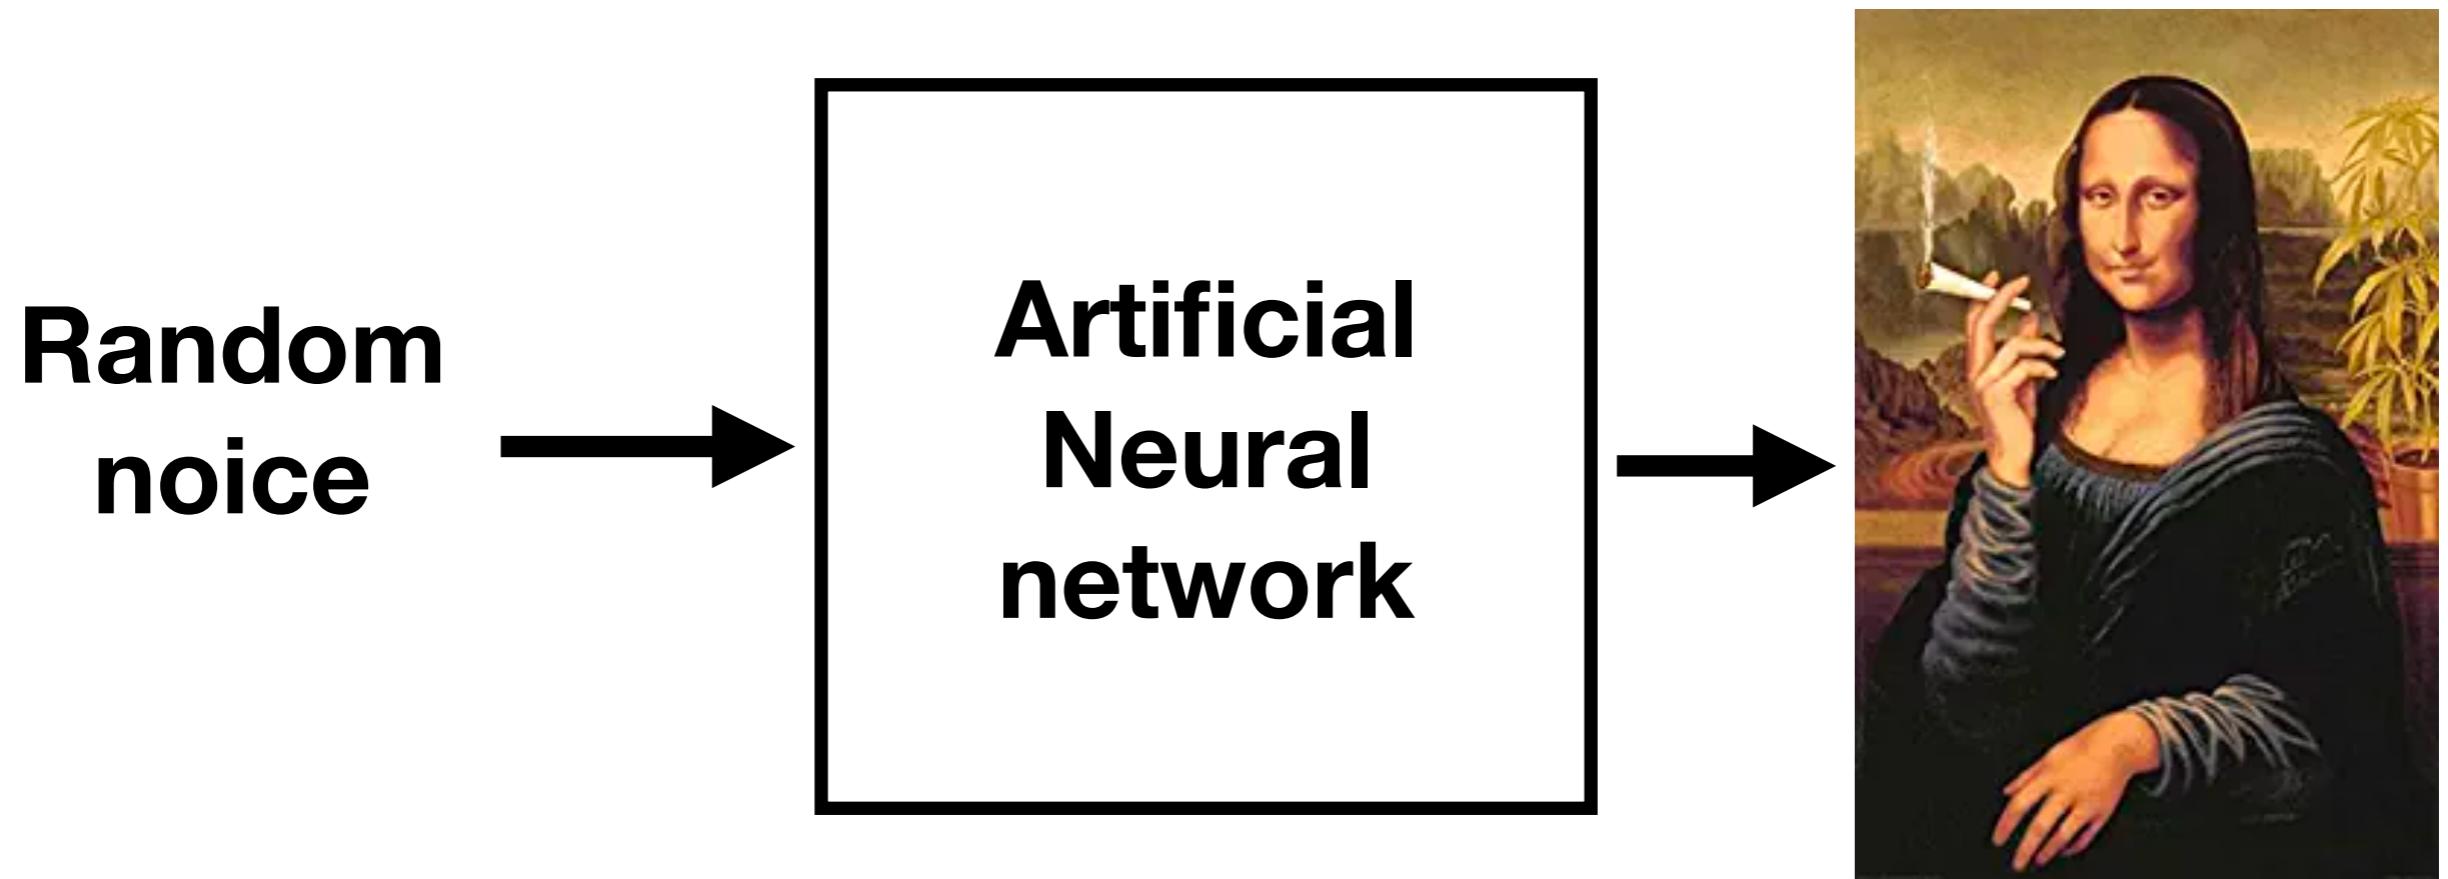
## Network Architecture



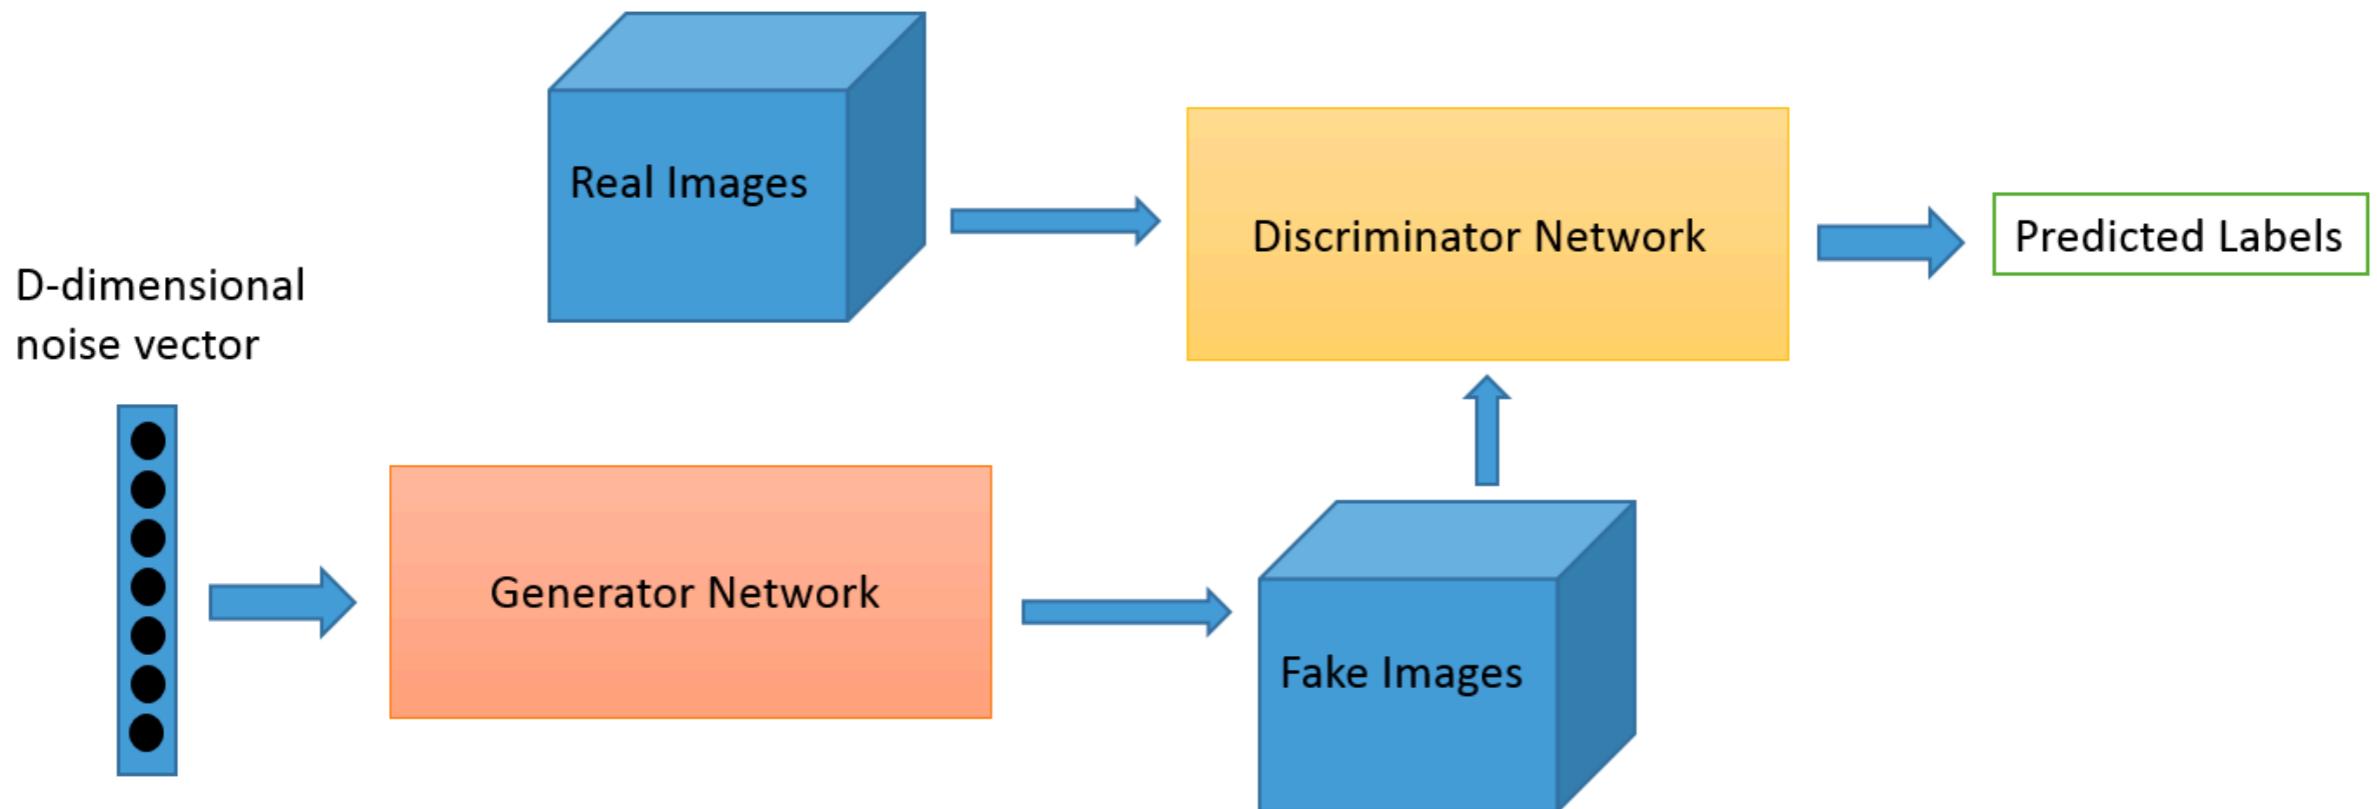
# **U-Net segmentation example**

**Segmentation.ipynb**

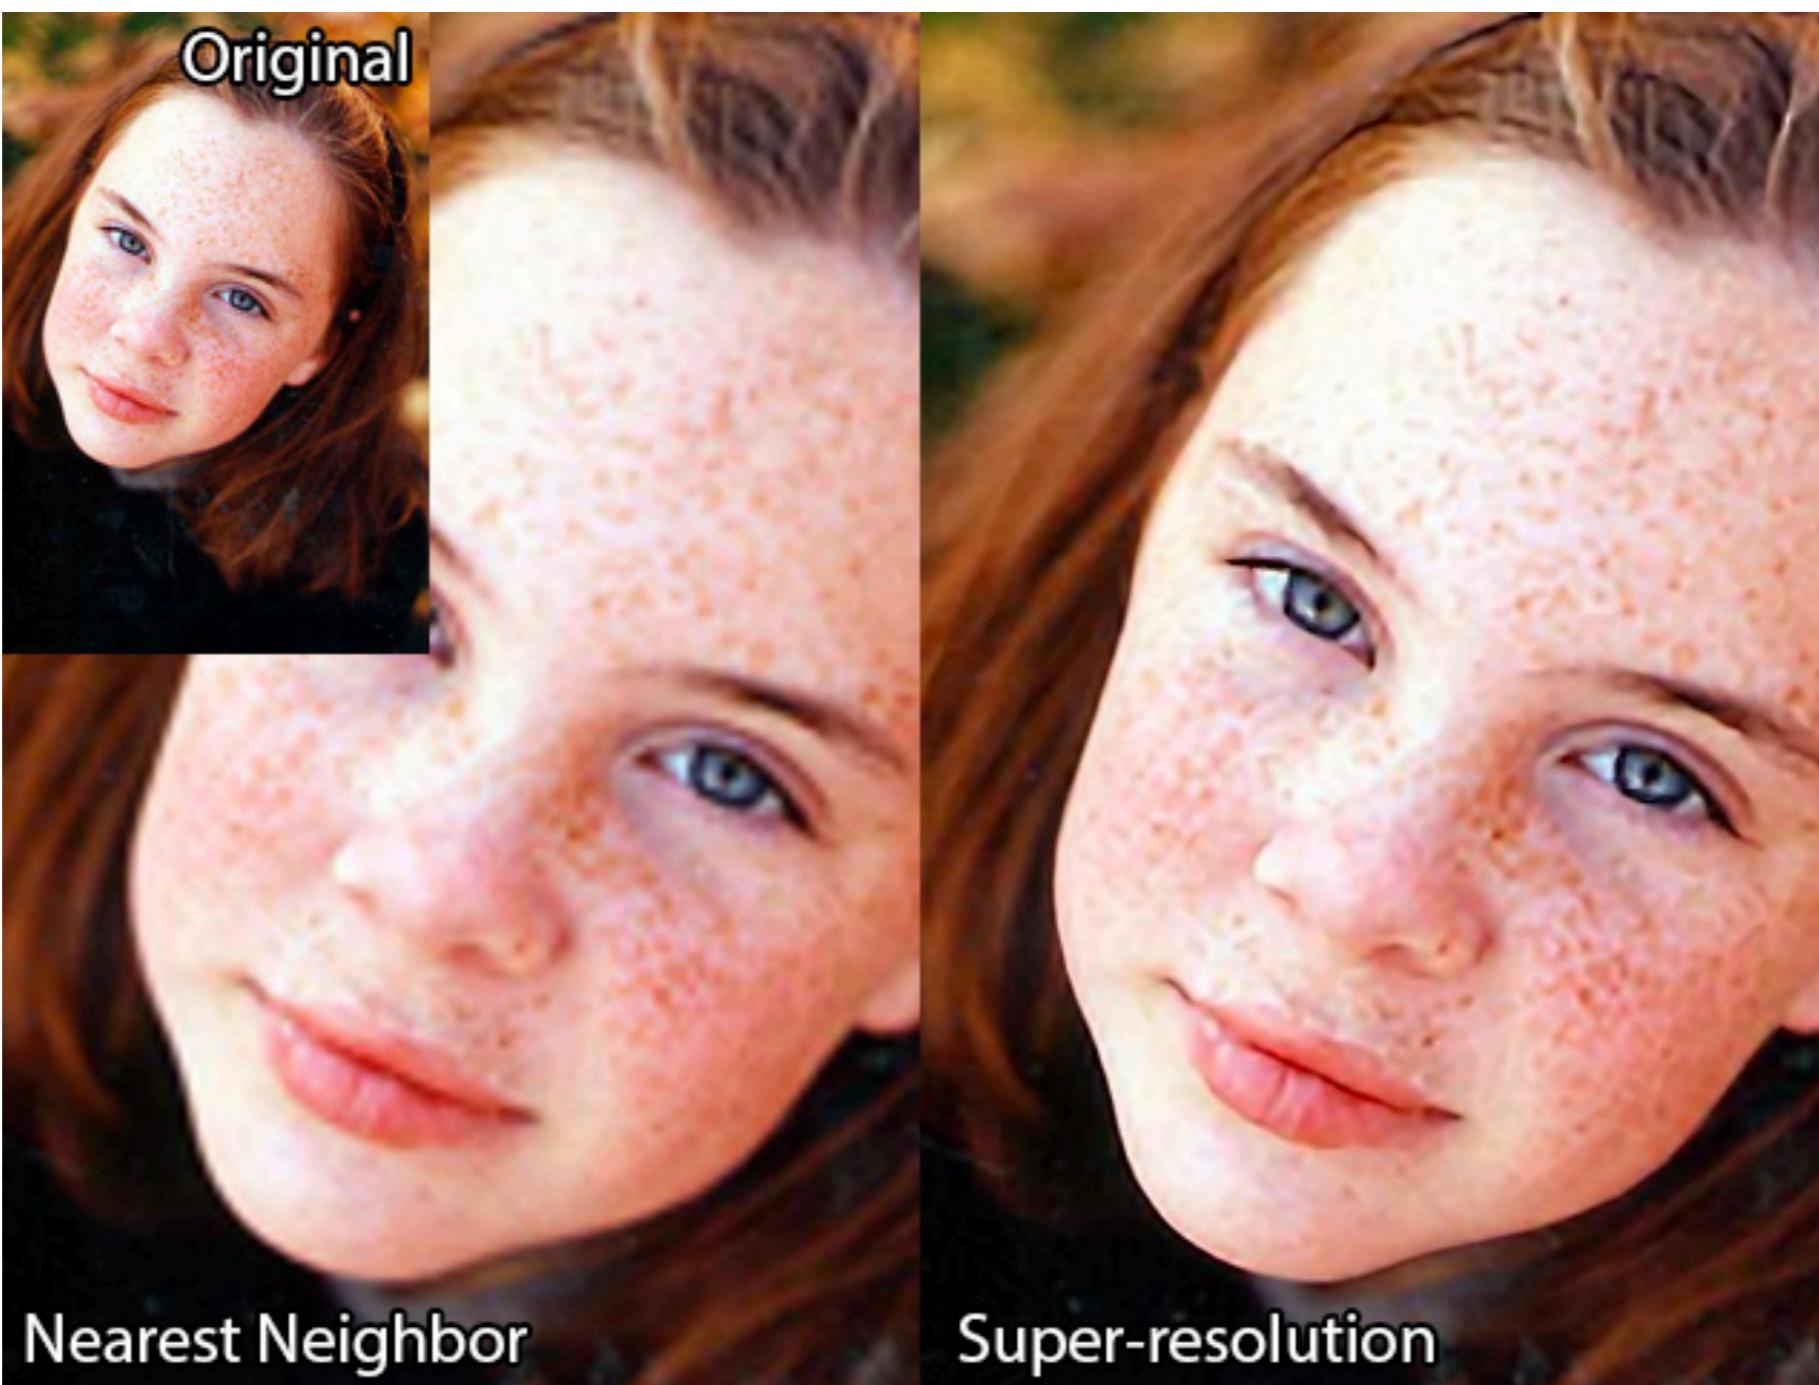
# Generative models with neural networks



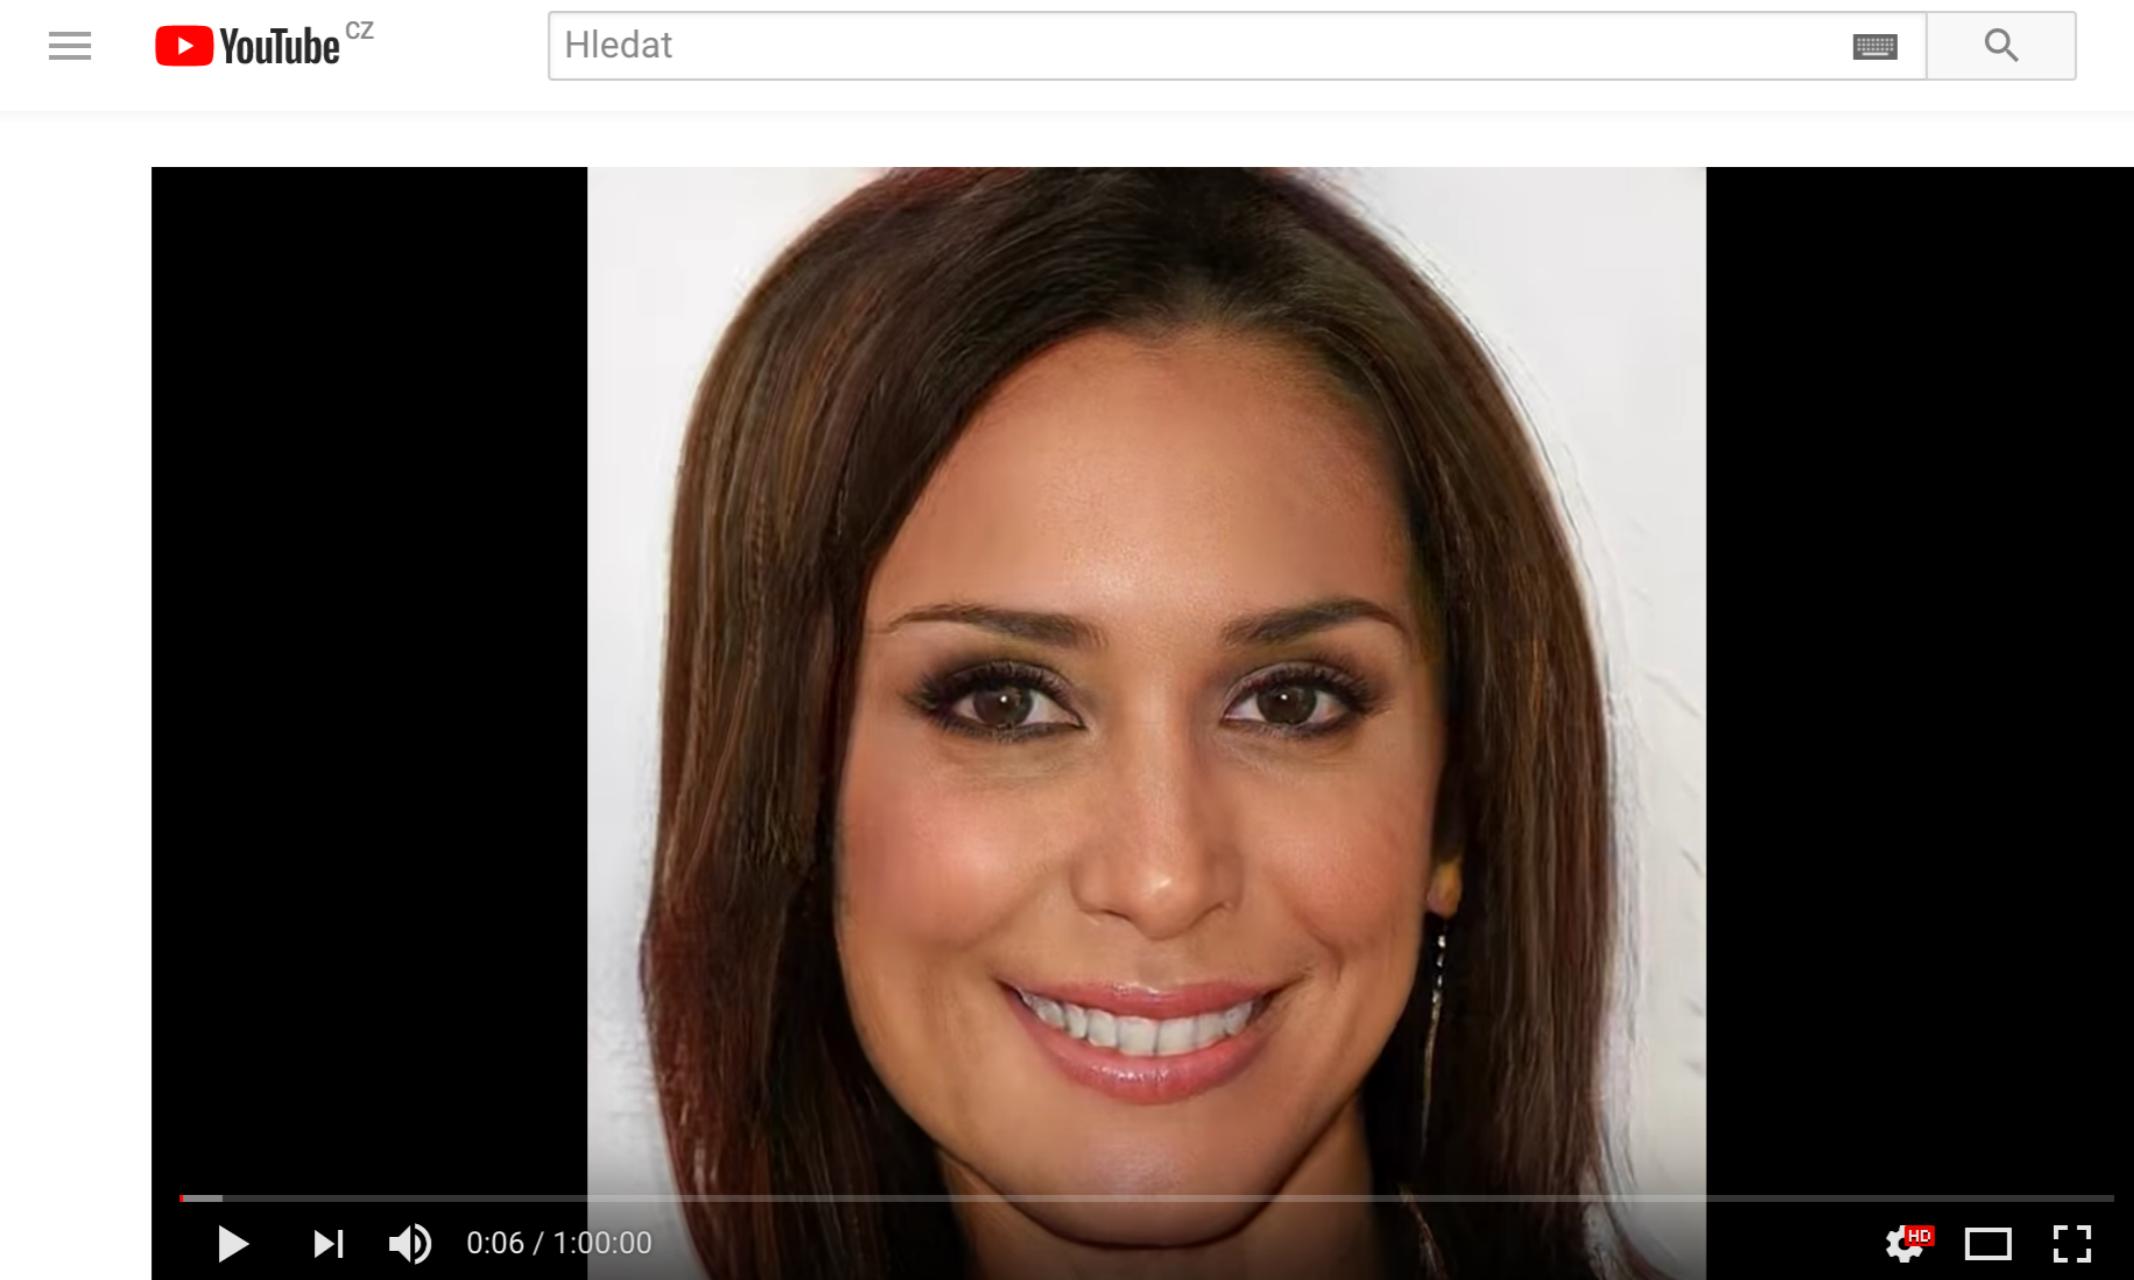
# Generative Adversarial Networks



# Superresolution



# Image synthesis



One hour of imaginary celebrities

95 832 zhlédnutí



TO SE MI LÍBÍ



NELÍBÍ SE



SDÍLET



...

# Which one is fake?



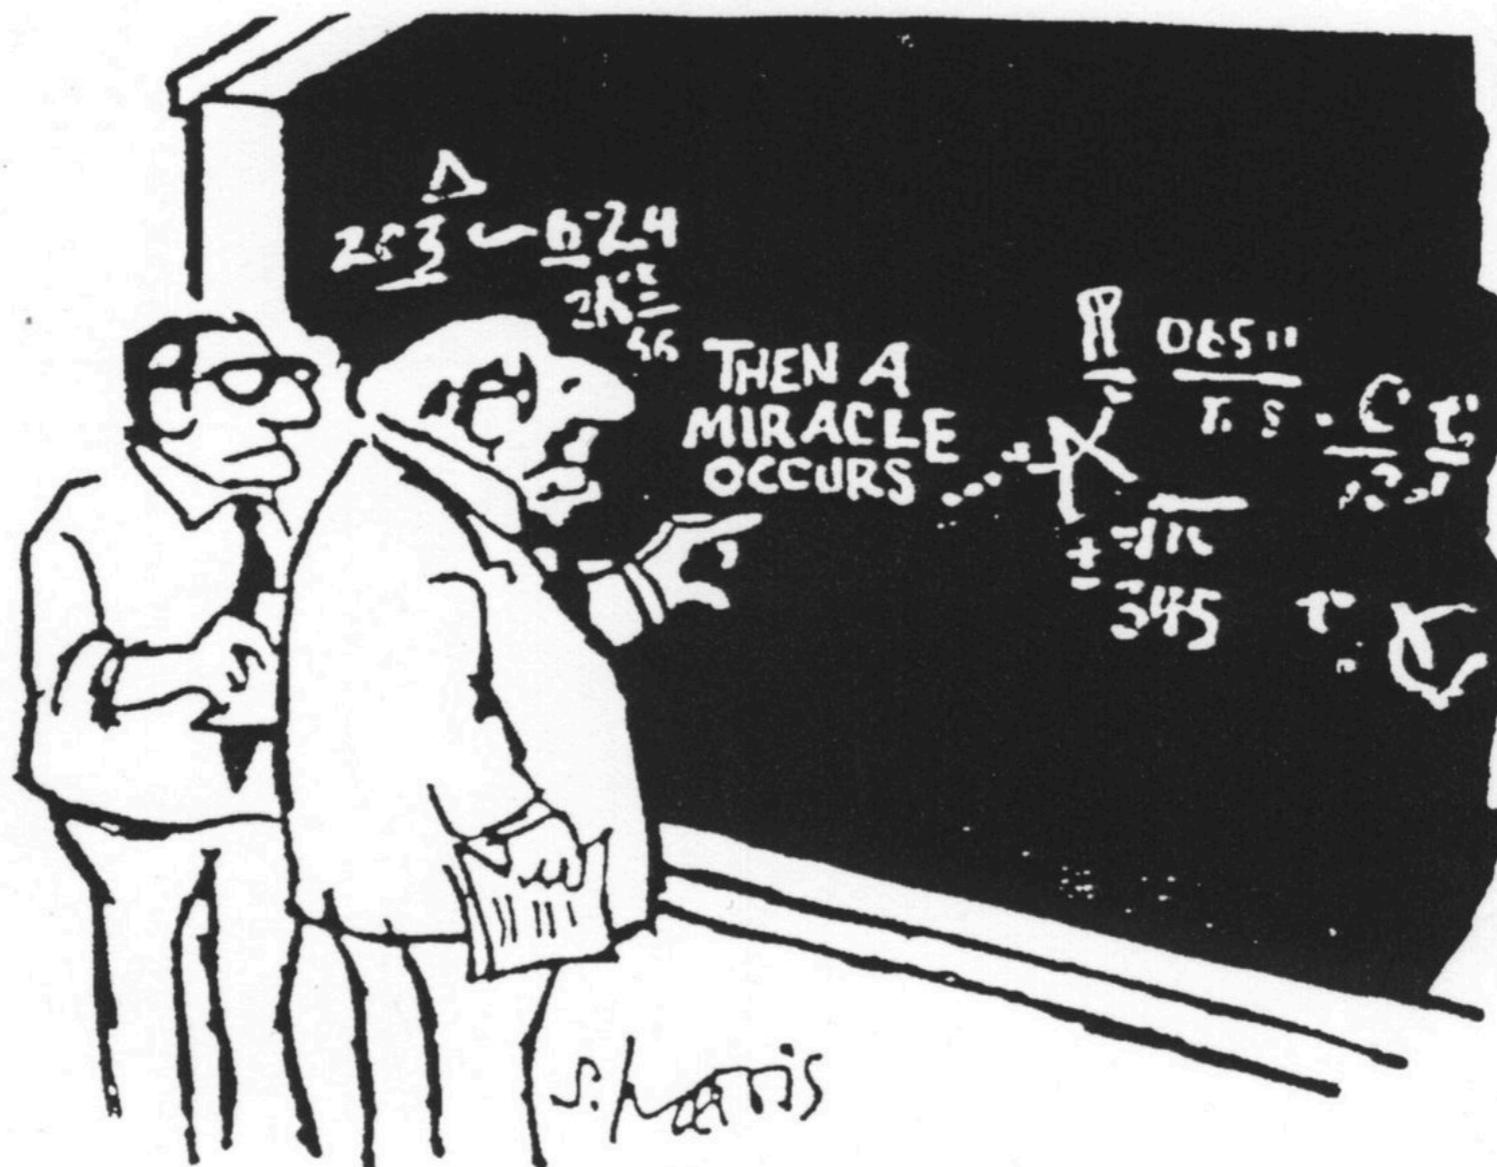
# Generative Adversarial Networks

**GANs.ipynb**

# Image manipulation



# Neural network explainability



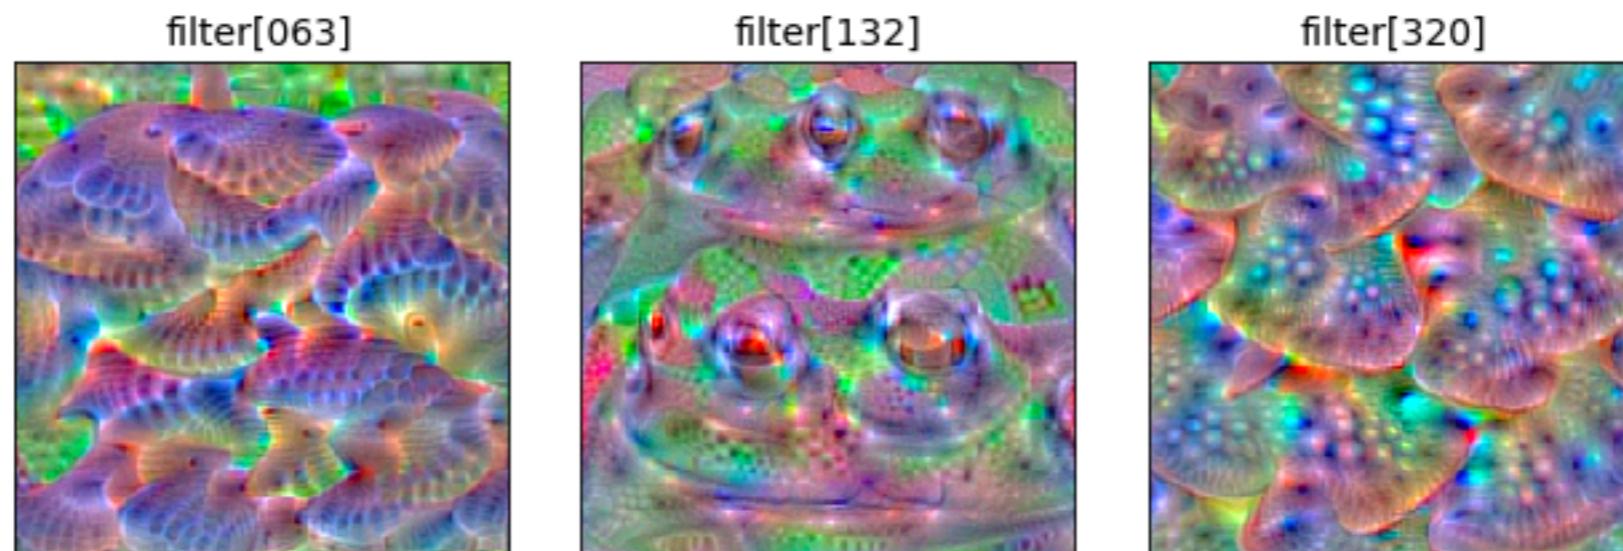
I think you should be a little  
more specific, here in Step 2

# Activation Maximization

## Visualized output classification Layer



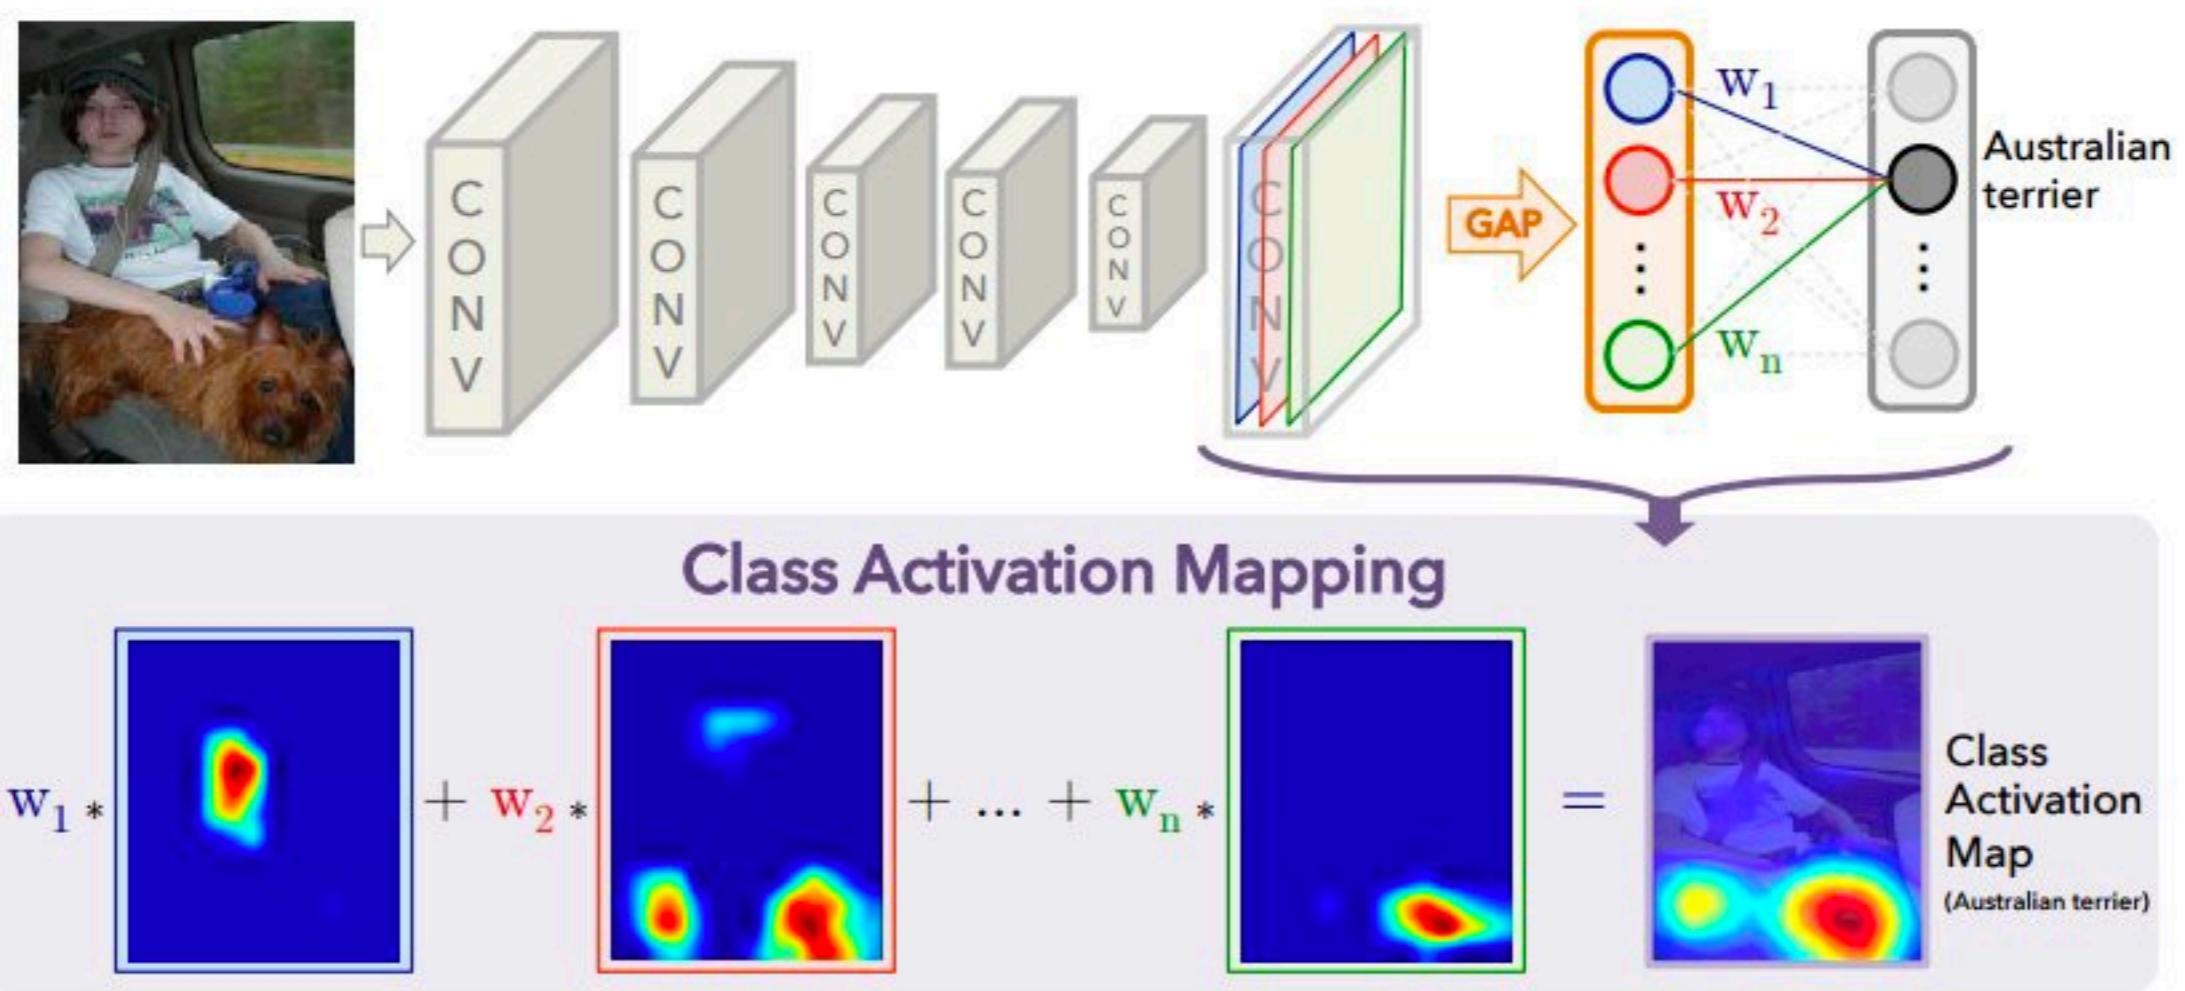
## Visualized hidden convolutional layers



# Grad-CAM heat maps



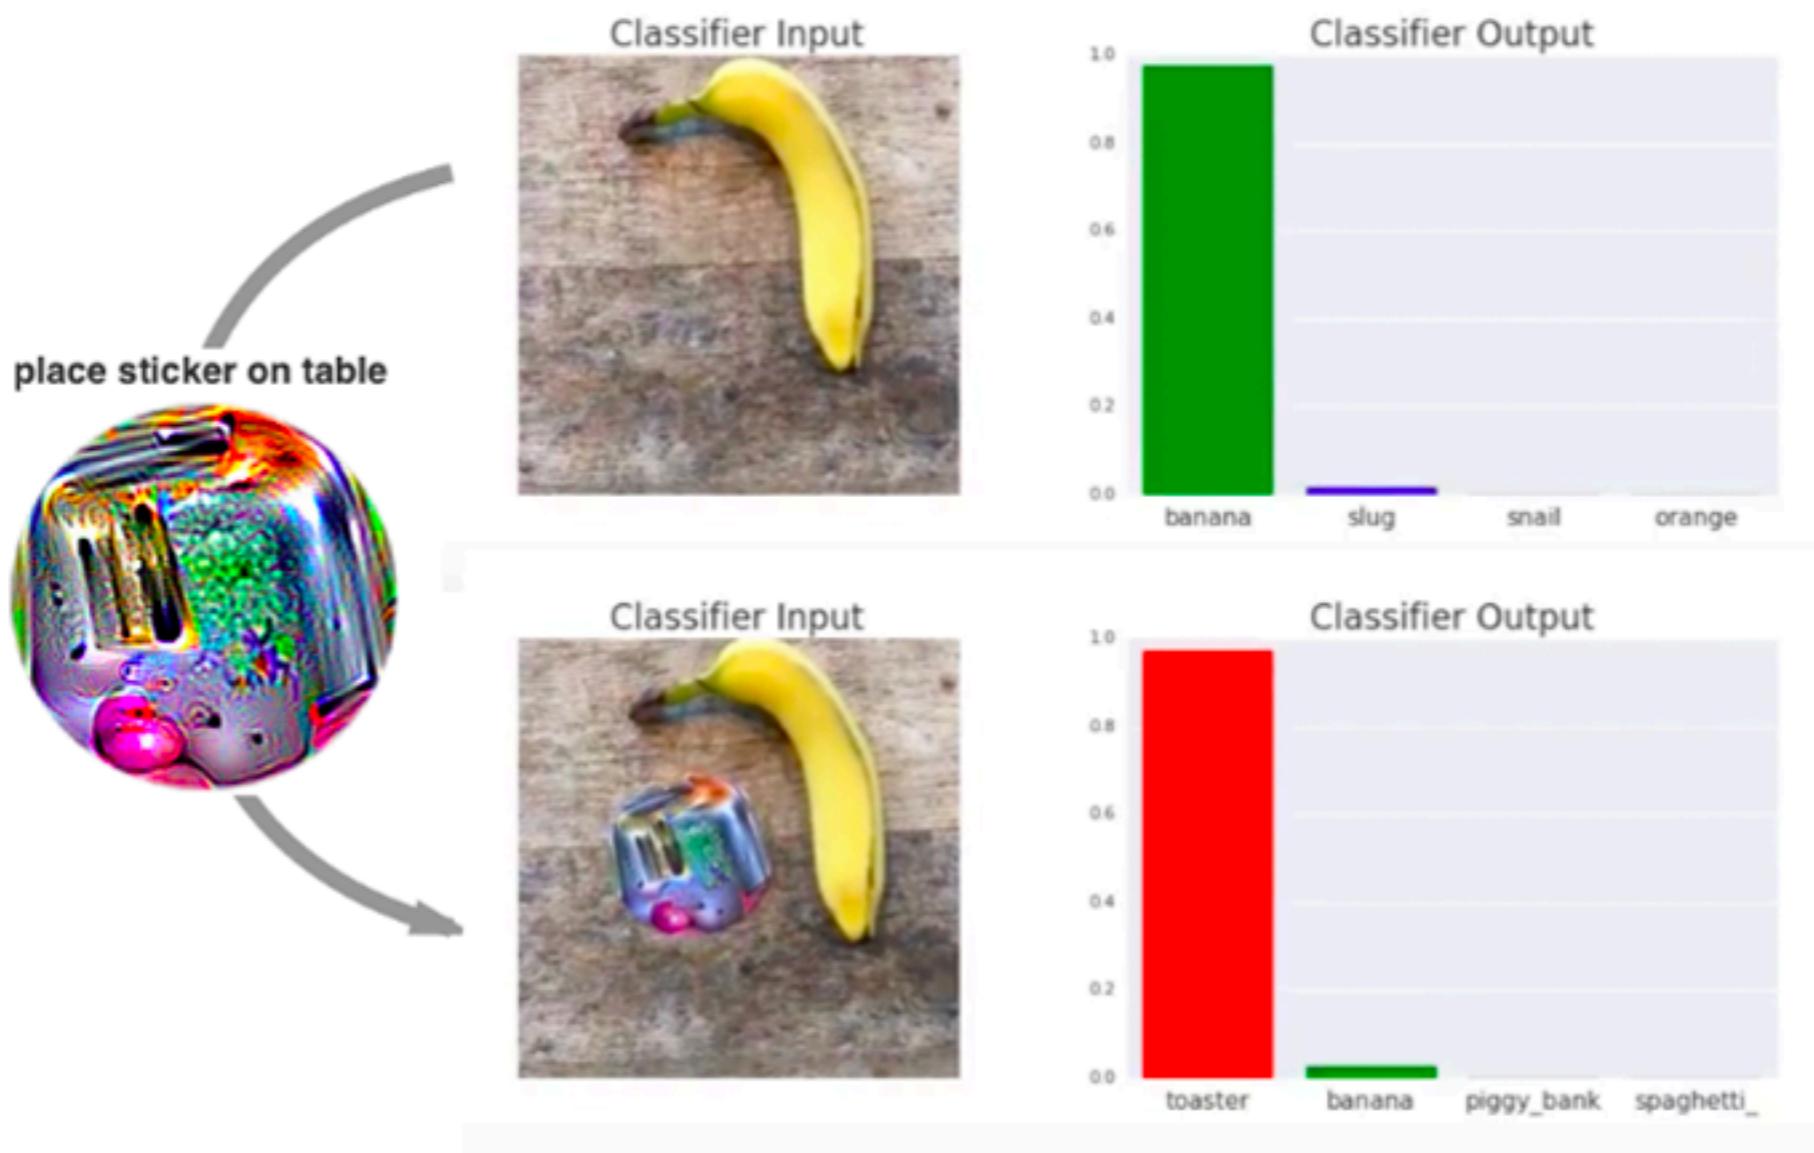
# CAM heat maps



# **Convolutional Neural Networks Explainability in Keras**

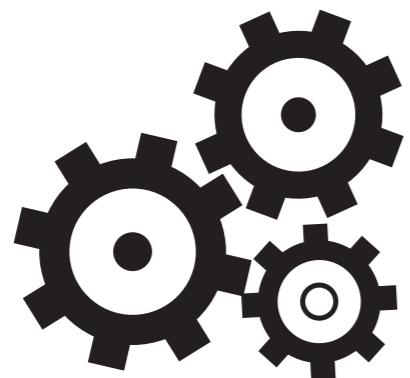
**Explainability.ipynb**

# Adversarial Patch



# What next?

**<https://www.mlcollege.com/en/#courses>**



Machine Learning Prague

**ML** MACHINE LEARNING  
**MU** meetups

# Thank you for your attention

**e-mail:** jiri@mlcollege.com

**Web:** www.mlcollege.com

**Twitter:** @JiriMaterna

**Facebook:** <https://www.facebook.com/maternajiri>

**LinkedIn:** <https://www.linkedin.com/in/jirimaterna/>